

GOLDEN GATE UNIVERSITY

Doctor of Business Administration (Generative AI)

DBA Dissertation v3

GGU-Strict Structure — Coverage-Matrix-Anchored

A Governance Framework for Responsible Explainable AI-Assisted Retrospective EEG-Based Neuropsychiatric Decision Support Under Human Clinical Oversight

*Integrating Generative AI, Retrieval-Augmented Generation,
and Continuous Monitoring for Clinical and Consumer Healthcare Applications*

Author: Praveen K Asthana

Cohort: GGU DBA 2026 • **Target Submission:** June 2026

Version: v3.0 • **Date:** Sunday 24th May, 2026 • **Pages:** 49

This dissertation v3 strictly mirrors the GGU Guidelines for Preparation of DBA Dissertation Thesis. Every GGU-required topic appears as a labelled section. A coverage matrix on page 2 maps every GGU bullet to its address.

Table of Contents (Topic Index)

1	GGU Coverage Matrix — Every Required Topic Addressed	7
2	Declaration	12
3	Certificate / Supervisor Approval	12
4	Acknowledgements	12
5	List of Abbreviations and Glossary of Terms	12
6	Abstract	14
7	Chapter 1 — Introduction	15
7.1	1.1 Background	15
7.2	1.2 Purpose	15
7.3	1.3 Problem Statement	15
7.4	1.4 Aim and Objectives	15
7.5	1.5 Research Questions	15
7.6	1.6 Hypothesis	16
7.7	1.7 Context of Contributing Organizations	16
7.8	1.8 Significant Contributions from the Investigation	16
7.9	1.9 Scope and Assumptions	17
7.10	1.10 Limitations	17
7.11	1.11 Thesis Outline	17
8	Chapter 2 — Review of Literature	18
8.1	2.1 Overview of Related Literature	18
8.2	2.2 Key Themes in the Literature	18
8.3	2.3 Strengths and Limitations of Previous Research	18
8.4	2.4 Identification of Research Gaps	18
8.5	2.5 Critical Analysis of Gaps and Challenges	19
8.6	2.6 Unresolved Issues in the Literature	19
8.7	2.7 Limitations of Existing Studies	19
8.8	2.8 Areas for Further Exploration	19
8.9	2.9 Summary of Key Insights from Literature	19
8.10	2.10 How the Study Contributes to Existing Knowledge	19
8.11	2.11 Justification for Research Approach	19
9	Chapter 3 — Research Methods	20
9.1	3.1 Research Strategy and Research Design	20
9.2	3.2 Research Approach (Qualitative, Quantitative, Mixed Methods)	21
9.3	3.3 Research Process	21
9.4	3.4 Data Sources (Primary and Secondary)	21
9.5	3.5 Data Collection Strategies	21
9.6	3.6 Sampling Strategies	21
9.7	3.7 Data Analysis Techniques	21

9.8	3.8 Limitations of Methodology	21
9.9	3.9 Ethical and Regulatory Considerations	21
9.10	3.10 Evaluation Metrics	22
9.11	3.11 Conclusion (Chapter 3)	22
10	Chapter 4 — Report on Data Analysis and Findings	23
10.1	4.1 Thematic Analysis of Collected Data	23
10.2	4.2 Key Findings	23
10.3	4.3 Quantitative and Qualitative Insights	23
10.4	4.4 Integration with Existing Systems	23
10.5	4.5 Findings in the Context	23
10.6	4.6 Contribution of the Study (Theoretical and Practical)	23
10.7	4.7 Ethical and Regulatory Considerations (Findings)	23
11	Chapter 5 — Discussion, Recommendations & Future Research	24
11.1	5.1 Interpretation of Findings	24
11.2	5.2 Implications for Business and Practice	24
11.3	5.3 Implications for Policy and Regulation	24
11.4	5.4 Limitations of the Study	24
11.5	5.5 Summary of Key Findings	24
11.6	5.6 Recommendations and Areas of Further Research	24
12	Appendices A–J — Planned Content	25
13	References (Selected Citations Aligned with Main Thesis)	25
14	Enhanced DBA Additions — Visuals, Tables, and Section Expansions	27
14.1	E1 — Figure: Conceptual Framework (5-Layer Pipeline)	27
14.2	E2 — Figure: Research Framework (Variables + Mediator-Moderator Paths)	27
14.3	E3 — Figure: RGAIG Enterprise Framework (5 Layers + Audit Fabric)	27
14.4	E4 — Figure: AS-IS vs TO-BE Swimlane (Caregiver Journey)	28
14.5	E5 — Figure: Sequence Diagram (Single Caregiver Query Through HITL Gate)	28
14.6	E6 — Table: Theme-Based Clustering of Literature	29
14.7	E7 — Table: Comparative Matrix (Study x Method x Industry x Limitation x Gap)	29
14.8	E8 — Table: Findings Traceability Matrix (Objective → RQ → Method → Analysis → Finding)	29
14.9	E9 — Table: Variables Definition (IV / Mediator / Moderator / DV + Qual/Quan + Measurement)	30
14.10	E10 — Table: Primary vs Secondary Data Split	30
14.11	E11 — Table: Theoretical vs Practical Contributions (Separated)	30
14.12	E12 — Table: Reliability and Validity (Cronbach α + Composite Reliability + HTMT)	31
14.13	E13 — Hypothesis Power Analysis (Expansion of §1.6)	31
14.14	E14 — Chapter Summaries (Ch. 2, Ch. 4, Ch. 5) + Final Conclusion	31
14.15	E15 — Hybrid (Mixed-Methods) Explicit Treatment	32
14.16	E16 — RGAIG Apply: How to Operationalise the Framework	32
14.17	E17 — Decision Analysis (Rule + Confidence + HITL Route)	33
14.18	E18 — Management Decision Matrix	33
14.19	E19 — Analysis Methods: Statistical + Subjective + Sensitivity (Monte Carlo Excluded)	33
14.20	E20 — Governance / AI / Compliance (Multi-Jurisdiction Crosswalk)	34

14.21E21 — 23-Step Complete Flow: ASD EEG → AI → RAG	34
14.22E22 — B2B and B2C Side-by-Side (Stakeholder Decomposition)	36
14.23E23 — Agentic AI Features: MCP + Council of Agents + Perplexity + Orchestration Patterns . .	36
14.24E24 — Production Infrastructure + Observability + CUA Stack	38
14.25E25 — Data Tier + Knowledge Layer (Storage + Streaming + Semantic Web)	40
14.26E26 — AI Security: Prompt Injection, OWASP LLM Top 10, Jailbreak Defence	41
14.27E27 — MLOps Lifecycle + 12-Tier Testing Matrix	41
14.28E28 — AI Discipline Map: 11 Cross-Cutting Concerns	42
14.29E29 — NIST AI RMF and ISO/IEC 42001 Detailed Compliance Mapping	43
14.30E30 — SMART Test of Research Aim and Objectives	45
14.31E31 — C4 Architecture Model (L1 System Context + L2 Containers + L3 Components + L4 Code)	46

List of Tables (Table Index)

1	GGU Coverage Matrix — 47 required topics, all addressed.	8
2	GGU ↔ Main Thesis Cross-Reference — every GGU bullet bridged to both the v3 section AND the substantive main-thesis location.	10
3	Acknowledgements — four constituencies + personal thanks.	12
4	Selected reference list (15 most-cited; full 424-entry bibliography in main thesis).	26
5	Six-theme clustering replacing author-by-author summary.	29
6	Comparative matrix skeleton (5 illustrative rows; dissertation expands to 20–30).	29
7	Findings traceability matrix — Objective → RQ → Method → Analysis → Finding.	29
8	Variables — IV / M1 / M2 / Moderator / DV with type + measurement.	30
9	Primary vs Secondary data split — 5 streams with IRB status.	30
10	Theoretical vs Practical contributions separated (per enhanced 4.10 / 4.11).	30
11	Reliability + validity tests with thresholds + reporting commitments.	31
12	Hypothesis-level power analysis (per H1–H5).	31
13	Chapter summaries + final conclusion (closing the enhanced structure gaps).	32
14	Hybrid (mixed-methods) design with both qual + quan legs + integration.	32
15	RGAIG Apply — 7-step operational playbook.	32
16	Decision-analysis routing — ML score → rule → confidence → HITL.	33
17	Management decision matrix — governance dimension → managerial decision.	33
18	Analysis methods — statistical + subjective + sensitivity + triangulation (Monte Carlo excluded per policy).	33
19	Governance / AI / Compliance crosswalk — RGAIG controls → 4 regulatory regimes.	34
20	23-step Neuropsychiatric EEG → AI → RAG pipeline with RGAIG layer + artefact per step.	35
21	B2B + B2C side-by-side decomposition — caregiver-primary vs clinician-supporting roles.	36
22	Agentic AI feature catalogue — A1–A10 with mitigated failure mode and audit-row field.	37
23	Agentic platform tool inventory — production-grade external tools and their RGAIG integration status.	38
24	Production infrastructure + observability + CUA stack with RGAIG context.	39
25	Data tier + knowledge layer — cache / streaming / RDBMS / vector / graph / RDF / OWL / ontology with RGAIG context.	40
26	AI security threat catalogue (OWASP LLM Top 10 + 3 beyond) with RGAIG mitigation per row.	41
27	MLOps lifecycle (M1–M8) + 12-tier testing matrix (T1–T12) with RGAIG gates.	42
28	AI discipline map — 11 cross-cutting concerns with RGAIG control + measurable threshold.	43
29	NIST AI RMF 1.0 — Govern / Map / Measure / Manage + cross-cutting characteristics mapped to RGAIG.	44
30	ISO/IEC 42001:2023 clauses + annexes mapped to RGAIG controls + audit evidence.	45
31	SMART aim — reformulated statement + per-criterion evidence.	46
32	SMART test of objectives O1–O6 — all six pass all five criteria with named evidence.	46
33	C4 L1 System Context — 4 actor types + 3 external system types + RGAIG boundary.	47
34	C4 L2 Containers — 8 deployable services + 4 persistent stores inside RGAIG boundary.	47
35	C4 L3 Components + L4 Code — reference pointers to main thesis Appendix I.	48

List of Figures (Figure Index)

1	Thesis outline — 5 chapters in logical sequence.	18
2	Research design — DSR 6-step loop with iteration.	20
3	Conceptual Framework — five layers from inputs to business value.	27
4	Research Framework — 1 IV + 2 Mediators + 1 Moderator + 1 DV.	27
5	RGAIG Enterprise Framework — 5 layers + transverse audit fabric.	28
6	AS-IS vs TO-BE swimlane — caregiver journey (5 steps each).	28
7	Sequence diagram — single caregiver query through HITL gate.	28
8	23-step pipeline in 7 phases + transverse governance + monitoring layer.	36
9	Agentic orchestration pattern — user → planner → decomposer → policy → MCP → Perplexity-RAG → council → HITL.	48

1. GGU Coverage Matrix — Every Required Topic Addressed

This matrix is the operator-facing contract: every bullet from the GGU "Guidelines for Preparation of DBA Dissertation Thesis" is named, located within this document, and certified as addressed. A green tick indicates the section is present with substantive content; the document is non-compliant if any row shows otherwise.

GGU #	GGU-required topic	Status	Addressed in section
Chapter 1 — Introduction (GGU target 20–30 pp; this v3 condenses to skeleton + content)			
1.1	Background	✓	§7.1
1.2	Purpose	✓	§7.2
1.3	Problem Statement	✓	§7.3
1.4	Aim and Objectives	✓	§7.4
1.5	Research Questions	✓	§7.5
1.6	Hypothesis	✓	§7.6
1.7	Context of Contributing Organizations (if required)	✓	§7.7
1.8	Significant contributions from the investigation	✓	§7.8
1.9	Scope and assumptions if any	✓	§7.9
1.10	Limitations	✓	§7.10
1.11	Thesis Outline	✓	§7.11
Chapter 2 — Review of Literature (GGU target 50–60 pp)			
2.1	Overview of Related Literature	✓	§8.1
2.2	Key Themes in the Literature	✓	§8.2
2.3	Strengths and Limitations of Previous Research	✓	§8.3
2.4	Identification of Research Gaps	✓	§8.4
2.5	Critical Analysis of Gaps and Challenges	✓	§8.5
2.6	Unresolved Issues in the Literature	✓	§8.6
2.7	Limitations of Existing Studies	✓	§8.7
2.8	Areas for Further Exploration	✓	§8.8
2.9	Summary of Key Insights from Literature	✓	§8.9
2.10	How the Study Contributes to Existing Knowledge	✓	§8.10
2.11	Justification for Research Approach	✓	§8.11
Chapter 3 — Research Methods (GGU target 40–50 pp)			
3.1	Research Strategy and Research Design	✓	§9.1
3.2	Research Approach (Qual / Quant / Mixed)	✓	§9.2
3.3	Research Process	✓	§9.3
3.4	Data Sources (Primary and Secondary)	✓	§9.4
3.5	Data Collection Strategies	✓	§9.5
3.6	Sampling Strategies	✓	§9.6
3.7	Data Analysis Techniques	✓	§9.7
3.8	Limitations of Methodology	✓	§9.8
3.9	Ethical and Regulatory Considerations	✓	§9.9
3.10	Evaluation Metrics	✓	§9.10
3.11	Conclusion (Ch.3)	✓	§9.11
Chapter 4 — Report on Data Analysis and Findings (GGU target 70–80 pp)			
4.1	Thematic Analysis of Collected Data	✓	§10.1
4.2	Key Findings	✓	§10.2
4.3	Quantitative and Qualitative Insights	✓	§10.3
4.4	Integration with Existing Systems	✓	§10.4
4.5	Findings in the context	✓	§10.5
4.6	Contribution of the study (Theoretical + Practical)	✓	§10.6
4.7	Ethical and Regulatory Considerations (findings)	✓	§10.7
Chapter 5 — Discussion, Recommendations & Future Research (GGU target 10–20 pp)			
5.1	Interpretation of Findings	✓	§11.1
5.2	Implications for Business and Practice	✓	§11.2
5.3	Implications for Policy and Regulation	✓	§11.3
5.4	Limitations of the Study	✓	§11.4
5.5	Summary of Key Findings	✓	§11.5
5.6	Recommendations and Areas of further research	✓	§11.6
Front matter + Appendices (Enhanced: recommended additions beyond strict GGU)			

Table 1: GGU Coverage Matrix — 47 required topics, all addressed.

GGU ⇔ Main-Thesis Cross-Reference Table

This second table is the bridge between this v3 GGU-strict structural index and the substantive main thesis (1126 pp); each GGU bullet maps to its v3 section AND to the exact main-thesis chapter / section that addresses it substantively. A professor can scan this single table to confirm every GGU requirement is addressed in both the structural map (v3) and the evidence-bearing main thesis.

GGU	GGU bullet	v3 section	Main-thesis chapter / section (where addressed substantively)
1.1	Background	§7.1	Ch. 1 §Background of the Study
1.2	Purpose	§7.2	Ch. 1 §Research Motivation + §Research Aim
1.3	Problem Statement	§7.3	Ch. 1 §Problem Statement (verbatim)
1.4	Aim and Objectives	§7.4	Ch. 1 §SMART Research Aim + §Objectives
1.5	Research Questions	§7.5	Ch. 1 §Research Questions
1.6	Hypothesis	§7.6	Ch. 4 §H1–H5 PLS-SEM tests (where evidence lives)
1.7	Contributing Organisations	§7.7	Ch. 1 §Context of Contributing Organizations
1.8	Significant contributions	§7.8	Ch. 1 §Significance of the Study + Ch. 5 §C1–C8
1.9	Scope and assumptions	§7.9	Ch. 1 §Scope and Boundaries of the Study
1.10	Limitations	§7.10	Ch. 1 §Delimitations + Ch. 5 §Limitations
1.11	Thesis Outline	§7.11	Ch. 1 §Organisation of the Thesis
2.1	Overview of Related Literature	§8.1	Ch. 2 §Introduction (opens with overview prose)
2.2	Key Themes	§8.2	Ch. 2 §Theoretical Framework (TAM + TOE + RBV + Trust)
2.3	Strengths and Limitations	§8.3	Ch. 2 throughout (per-theme critique)
2.4	Research Gaps	§8.4	Ch. 2 §Research Gaps G1–G6
2.5	Critical Analysis of Gaps	§8.5	Ch. 2 §Critical Analysis (per gap)
2.6	Unresolved Issues	§8.6	Ch. 2 §Open Issues + Ch. 5 §Future Research
2.7	Limitations of Existing Studies	§8.7	Ch. 2 §Limitations of Prior Work
2.8	Areas for Further Exploration	§8.8	Ch. 5 §Future Research Areas
2.9	Summary of Key Insights	§8.9	Ch. 2 §Chapter Summary
2.10	Contribution to Existing Knowledge	§8.10	Ch. 5 §Theoretical Contributions C1–C8
2.11	Justification for Research Approach	§8.11	Ch. 3 §Research Philosophy + §Strategy
3.1	Research Strategy and Design	§9.1	Ch. 3 §Research Strategy + §Research Design
3.2	Research Approach	§9.2	Ch. 3 §Mixed-Methods Approach
3.3	Research Process	§9.3	Ch. 3 §Research Process (DSR loop)
3.4	Data Sources	§9.4	Ch. 3 §Data Strategy (Streams 1–5)
3.5	Data Collection	§9.5	Ch. 3 §Prior IRB Approval Trail + Data Collection
3.6	Sampling	§9.6	Ch. 3 §Sampling Strategy
3.7	Data Analysis	§9.7	Ch. 3 §Statistical Methods + Ch. 4 §PLS-SEM
3.8	Methodological Limitations	§9.8	Ch. 3 §Limitations + Ch. 5 §Limitations
3.9	Ethics and Regulatory	§9.9	Ch. 3 §Ethical Considerations + App. D
3.10	Evaluation Metrics	§9.10	Ch. 3 §Evaluation Metrics + Ch. 4 §Findings
3.11	Conclusion (Ch. 3)	§9.11	Ch. 3 §Chapter Summary
4.1	Thematic Analysis	§10.1	Ch. 4 §Qualitative Findings + App. E coding maps
4.2	Key Findings	§10.2	Ch. 4 §H1–H5 Results + summary table
4.3	Quan + Qual Insights	§10.3	Ch. 4 §Triangulation
4.4	Integration with Existing Systems	§10.4	Ch. 4 §Workflow Integration + App. J
4.5	Findings in Context	§10.5	Ch. 4 §Discussion of Findings
4.6	Theoretical + Practical Contributions	§10.6	Ch. 4 §Contributions + Ch. 5 §C1–C8
4.7	Ethics + Regulatory (findings)	§10.7	Ch. 4 §Fairness + Governance Findings
5.1	Interpretation of Findings	§11.1	Ch. 5 §Interpretation of Findings
5.2	Business and Practice Implications	§11.2	Ch. 5 §Business Implications + ROI
5.3	Policy and Regulatory Implications	§11.3	Ch. 5 §Policy Implications + multi-jurisdiction
5.4	Limitations of the Study	§11.4	Ch. 5 §Limitations (full restatement)
5.5	Summary of Key Findings	§11.5	Ch. 5 §Summary of Findings
5.6	Recommendations + Further Research	§11.6	Ch. 5 §Recommendations + §Future Research

Table 2: GGU ↔ Main Thesis Cross-Reference — every GGU bullet bridged to both the v3 section AND the substantive main-thesis location.

2. Declaration

The candidate declares that this dissertation is original work and that all AI tool usage is disclosed in Appendix I per institutional policy. Signature blocks remain blank for institutional signing and are a release blocker before final submission.

Declaration item	Statement
Originality	Original work; no fabrication; no plagiarism.
AI tool disclosure	Codex + Claude used for editorial / structuring assistance; all decisions verified by candidate (Appendix I).
Prior IRB approval	Streams 3 + 4 ($n=131 + n=12$) cited as aggregated, anonymised, prior-IRB-approved secondary data.
Submission readiness	Topic proposal → dissertation alignment verified (separate audit doc).

3. Certificate / Supervisor Approval

This certificate carries the supervisor's confirmation that the topic, scope, and methodology are appropriate for the DBA submission. Signature blocks are left blank for institutional signing.

Role	Name (insert)	Signature / Date
Doctoral Supervisor	C. Ranjeeth Kumar	
Co-Supervisor / Chair	Sumitra Padmanabhan	
DBA Programme Director	<i>[To be filled by GGU Programme Office at submission]</i>	

4. Acknowledgements

The acknowledgements page recognises the people and institutions whose support made this doctoral journey possible. The candidate offers thanks across four constituencies — academic, institutional, data-source, and personal — each named in the table below.

Constituency	Acknowledgement
Doctoral supervision	Sincere thanks to C. Ranjeeth Kumar (Doctoral Supervisor) for guiding the research direction, framing the RGAIG construct, and challenging the work toward DBA-grade rigour.
Co-supervision / Chair	Sincere thanks to Sumitra Padmanabhan (Co-Supervisor / Chair) for the chairing oversight, methodological critique, and committee facilitation throughout the dissertation cycle.
GGU School of Business	Thanks to the Doctorate in Business Administration (Generative AI) cohort, the DBA Programme Office, and the doctoral peer-review group at Golden Gate University, San Francisco.
Data custodians	Acknowledgement to the King Abdulaziz University (KAU) team for the publicly available pediatric ASD-EEG dataset and the ABIDE-II consortium for the multi-site reference dataset; both were essential to the retrospective evidence chapter.
Prior research contributors	The candidate acknowledges the contributors to the prior-IRB-approved 131-clinician PLS-SEM cohort and the 12-expert qualitative panel whose aggregated, anonymised data underpins the empirical analysis in this dissertation.
Personal	Heartfelt thanks to family for the patience, encouragement, and resilience that accompanied this multi-year doctoral effort, and to the broader community of practitioners who shared the burdens and aspirations that motivate this work.

Table 3: Acknowledgements — four constituencies + personal thanks.

5. List of Abbreviations and Glossary of Terms

The abbreviation and glossary table consolidates the multi-disciplinary vocabulary spanning AI engineering, governance, healthcare, and DBA practice. A single canonical list prevents inconsistency across chapters and supports examiner navigation.

Abbrev.	Expansion / definition
RGAIG	Responsible Generative AI Governance — the proposed framework.
ASD	Autism Spectrum Disorder.
EEG	Electroencephalogram.
HITL	Human-In-The-Loop — clinician oversight gate.
XAI	Explainable AI (SHAP, LIME, counterfactual).
RAI	Responsible AI (fairness + accountability + transparency + privacy).
RAG	Retrieval-Augmented Generation.
DSR	Design Science Research.
PLS-SEM	Partial Least Squares Structural Equation Modelling.
TAM / TOE / RBV	Technology Acceptance / Tech-Org-Environment / Resource-Based View.
KAU / ABIDE-II	Saudi reference EEG dataset / Western multi-site ASD dataset.
NIST AI RMF	US National AI Risk Management Framework.
ISO/IEC 42001	International AI Management System standard.
EU AI Act	European Union Artificial Intelligence Act.
HIPAA / GDPR / DPDP	US health-data / EU general data / India data-protection.

6. Abstract

The abstract distils the dissertation into seven structured movements following the enhanced DBA convention. Examiners read the abstract first; the table below mirrors the seven-part skeleton mandated by the enhanced DBA structure.

Abstract movement	Content of this dissertation
1. Research Background	ASD prevalence rising; pediatric screening burden in tier-1/tier-2 institutions; Indian under-representation; growing Generative AI capability with immature governance fabric.
2. Research Problem	Enterprise / clinical AI adoption fails trust, explainability, accountability, and adoption gates simultaneously. Current frameworks address these in isolation.
3. Aim and Objectives	Develop and evaluate the RGAIG framework for B2C-primary, B2B-supporting AI-assisted ASD screening support; 6 objectives, 8 RQs, 5 hypotheses.
4. Methodology	Pragmatist mixed-methods Design Science Research + PLS-SEM ($n=131$) + qualitative expert panel ($n=12$) + retrospective EEG benchmarks (KAU + ABIDE-II).
5. Key Findings	H1–H5 evidence that governance maturity mediates between RAI principles and adoption; explainability + emotional reassurance dominate caregiver trust.
6. Contribution	Theoretical (RGAIG construct + 525-module taxonomy), practical (maturity instrument), methodological (DSR + PLS-SEM hybrid), policy (multi-jurisdiction governance mapping).
7. Keywords	Generative AI; Responsible AI; AI Governance; Enterprise AI; Agentic AI; Explainable AI; Digital Transformation; ASD; EEG; Healthcare AI.

7. Chapter 1 — Introduction

7.1 1.1 Background

This subsection establishes industry evolution, AI evolution, and the digital-transformation context that motivates the research, drawing on the AI-ethics critique of *principles alone* [1] and the healthcare-AI assurance literature [2]. The dissertation tracks adoption statistics, business pain points, and the regulatory timeline (NIST AI RMF [11], ISO/IEC 42001 [12], EU AI Act [14], DPDP [15]) that creates the window for a governance contribution.

Background dimension	What the dissertation establishes
Industry evolution	Pediatric ASD prevalence rising globally; specialist wait 3–12 months; tier-2/3 city access constrained.
AI evolution	GenAI + RAG + agentic systems mature 2023–2026; immature governance fabric.
DT context	Healthcare DT stalled by trust, explainability, governance gaps.
Regulatory timing	EU AI Act + ISO/IEC 42001 + NIST AI RMF + DPDP all live in 2024–2026.
Market trends	Multi-jurisdiction governance harmonisation creates the contribution window.

7.2 1.2 Purpose

The purpose of this dissertation is to develop and evaluate a Responsible Generative AI Governance framework (RGAIG) that closes the integrated trust + explainability + adoption gap for pediatric ASD screening support. The purpose is non-diagnostic and positioned as B2C-primary, B2B-supporting per the dissertation's scope contract.

7.3 1.3 Problem Statement

Pediatric ASD screening support is fragmented, slow, and untrustworthy at the caregiver–clinician interface [10]; current AI tools either ignore the caregiver (B2B-only) or ignore the clinician's authority (B2C-only), with measurable economic burden [11] and emotional cost on families [14]. The five-element problem decomposition below is the load-bearing claim the dissertation defends in evidence chapters.

Element	Statement
Current state	Specialist wait 3–12 mo; black-box outputs; jargon reports; no SOS channel.
Existing limitation	Frameworks address governance OR explainability OR trust OR adoption — none address all four.
Business impact	Healthcare AI adoption stalls at procurement; trust erodes on first surprise.
Research gap	Integrated governance + explainability + trust + adoption fabric for pediatric AI.
Why current solutions fail	Single-axis design; no clinician HITL gate; no caregiver-readable XAI; no multi-jurisdiction mapping.

7.4 1.4 Aim and Objectives

The aim is one sentence; the six objectives are the executable pieces of that aim and are each testable, time-bounded, and traceable to a chapter. This decomposition mirrors the enhanced DBA convention for high-impact objectives.

#	Aim / Objective
Aim	Develop and evaluate the RGAIG framework for B2C-primary, clinician-supported AI-assisted ASD screening support.
O1	Investigate governance gaps in pediatric AI screening support (Ch. 1, Ch. 2).
O2	Evaluate existing frameworks (NIST AI RMF, ISO/IEC 42001, EU AI Act, DPDP) (Ch. 2, Ch. 3).
O3	Design the RGAIG framework artefact (5 layers + audit fabric) (Ch. 3, Ch. 5).
O4	Validate framework effectiveness empirically (PLS-SEM + panel + benchmark) (Ch. 4).
O5	Assess business + operational + ROI implications (Ch. 5).
O6	Articulate the adoption pathway + maturity instrument (Ch. 4, Ch. 5).

7.5 1.5 Research Questions

Eight research questions decompose the aim into focused, falsifiable inquiries; each RQ pairs with a hypothesis (where quantitative) and traces to a chapter that answers it. A DBA without a numbered RQ block is treated as exploratory; the numbered list below is the dissertation's commitment to confirmatory contribution.

RQ	Question	Answered in
RQ1	What governance gaps exist in pediatric AI screening support?	Ch. 2 + Ch. 4 thematic.
RQ2	How do RAI principles improve trust?	Ch. 4 PLS-SEM (H1, H2).
RQ3	Which operational controls reduce AI risk?	Ch. 3 + Ch. 4.
RQ4	Which mediators explain adoption?	Ch. 4 PLS-SEM (H3, H4).
RQ5	How does HITL effectiveness moderate the trust → adoption path?	Ch. 4 PLS-SEM (H5).
RQ6	What is the integration overhead of RGAIG in clinical workflow?	Ch. 4 §4.8.
RQ7	How does RGAIG perform on retrospective KAU + ABIDE-II benchmarks?	Ch. 4 §4.7.
RQ8	What policy + regulatory implications follow for multi-jurisdiction deployment?	Ch. 5 §5.3.

7.6 1.6 Hypothesis

Five hypotheses operationalise the conceptual framework into testable claims; each pairs with a null counterpart to satisfy Popperian falsifiability. PLS-SEM in Ch. 4 §4.5 tests every H1–H5 path with bootstrap mediation analysis.

H	Hypothesis + null
H1	RAI principles positively influence governance maturity. (H1 ₀ : no effect.)
H2	Governance maturity positively influences trust + adoption + adherence. (H2 ₀ : no effect.)
H3	Governance maturity <i>mediates</i> RAI → adoption. (H3 ₀ : no mediation.)
H4	Explainability fit <i>mediates</i> RAI → trust (partial). (H4 ₀ : no mediation.)
H5	HITL effectiveness <i>moderates</i> mediator → adoption. (H5 ₀ : no moderation.)

7.7 1.7 Context of Contributing Organizations

Per GGU guideline this section is conditional; the dissertation operates with prior-IRB-approved secondary data and publicly available datasets, so no operational contributing organisation is on the critical path of the current submission. The table below documents which organisations are referenced contextually and which would be engaged for Phase 2.

Organisation / dataset	Role	Status in current submission
KAU EEG (Saudi Arabia)	Reference dataset	Used (public, secondary).
ABIDE-II consortium (US/EU)	Reference dataset	Used (public, secondary).
GGU School of Business	Awarding body	Supervisor + chair sign-off pending.
Indian pediatric institutions (Phase 2)	Future cohort partner	EXCLUDED from current IRB; separate amendment planned.

7.8 1.8 Significant Contributions from the Investigation

Eight contributions C1–C8 span theoretical, practical, methodological, and policy dimensions; each is testable and traceable to a chapter. The dissertation defends each contribution against the prior-art comparative matrix in Ch. 2 §2.6.

C#	Type	Contribution
C1	Theoretical	RGAIG construct integrating RAI + governance + trust + adoption.
C2	Theoretical	525-module unified quality taxonomy (8-phase ML + 13-phase GenAI QA + 10 pillars).
C3	Methodological	DSR + PLS-SEM + qualitative-panel hybrid evaluation design.
C4	Practical	RGAIG maturity instrument (6-level ladder, 30-item scale).
C5	Practical	24-control catalogue mapped to NIST + ISO + EU AI Act + DPDP.
C6	Practical	Caregiver-readable + clinician-grade two-layer XAI specification.
C7	Policy	Multi-jurisdiction governance crosswalk (4 regulatory regimes).
C8	Practical	7-step caregiver adoption journey + 5-phase trust-building journey.

7.9 1.9 Scope and Assumptions

The scope is explicitly bounded along five dimensions; the assumptions are stated openly so reviewers can evaluate generalisability claims against them. A scope that cannot be stated in one table is too ambiguous for examiner defence.

Scope / assumption	Statement
Geography	Reference data Saudi Arabia (KAU) + multi-site Western (ABIDE-II); Indian cohort Phase 2 future.
Industry	Pediatric ASD screening support; healthcare AI; non-diagnostic.
Technology boundary	Multi-agent + RAG + HITL; reference implementation in candidate's <code>agenticfinder</code> codebase.
Time boundary	2024–2026 regulatory texts; technology snapshot June 2026.
Organisational scope	GGU as awarding body; no operational deployment in current IRB.
Assumption 1	Baseline digital infrastructure exists at adopting sites.
Assumption 2	Clinician availability for HITL gate exists at adopting sites.
Assumption 3	Participant + data-source AI awareness assumed.

7.10 1.10 Limitations

The dissertation declares six limitations explicitly so an examiner can interrogate each one in turn. Limitations are not weaknesses if they are named, scoped, and bounded by mitigations.

Limitation	Mitigation in dissertation
Sample size (n=131 quan; n=12 qual)	Justified by bootstrap power analysis; transparent CI reporting.
Retrospective bias	Explicit non-claim of prospective deployment; scope contract in §1.9.
Cross-cultural generalisability	Indian-cohort Phase 2 in Appendix N (excluded from current IRB).
Technology evolution	Model + prompt versioning; audit row carries <code>model_v</code> + <code>prompt_v</code> .
Access restrictions	Phase 2 outreach pack + tracker; multi-institution discovery framework.
Regulatory evolution	Multi-jurisdiction crosswalk versioned; texts cited dated 2025–2026.

7.11 1.11 Thesis Outline

The thesis flows through 5 chapters whose outputs feed each other: Ch. 1 introduces, Ch. 2 grounds, Ch. 3 designs, Ch. 4 evidences, Ch. 5 interprets and prescribes. The flowchart below is the operator-facing one-page chapter map per GGU guideline.




Figure 1: Thesis outline — 5 chapters in logical sequence.

8. Chapter 2 — Review of Literature

8.1 2.1 Overview of Related Literature

The literature spans AI governance [7], explainable AI [6], responsible AI [1], agentic systems [12], RAG, pediatric ASD diagnostics [10], healthcare AI [2], and digital-transformation theory anchored in TAM [3] and trust theory [4, 5]. The dissertation organises this multi-axis literature via theme-based clustering rather than author-by-author summary so contradictions and gaps surface visibly.

8.2 2.2 Key Themes in the Literature

Six themes cluster the literature into examinable categories; each theme is paired with its focus and the representative work the dissertation engages. Theme-based clustering is the high-impact alternative to chronological summary mandated by enhanced DBA practice.

Theme	Focus	Representative work + open issue
AI Governance	Risk + compliance	NIST AI RMF; ISO/IEC 42001; EU AI Act — integration with operational HITL underspecified.
Explainability	Transparency	SHAP, LIME, counterfactual — caregiver vs. clinician layer distinction missing.
Agentic systems	Autonomous reasoning	Planner-decomposer-policy architectures — deterministic auditability open.
RAG architectures	Retrieval optimisation	Naive → advanced → agentic-RAG — hallucination control at long context open.
Human-in-the-Loop	Oversight	HITL gating + escalation — cost-benefit at scale open.
AI Ethics	Fairness + accountability	Mittelstadt principles; EU AI Act Art. 86 — pediatric-specific safeguards open.

8.3 2.3 Strengths and Limitations of Previous Research

Per theme, the dissertation identifies strengths the literature has earned and limitations it has not yet resolved. This two-column framing makes the literature review evaluative rather than descriptive.

Theme	Strengths in prior work	Limitations / unresolved
AI Governance	Mature compliance frameworks	Operational HITL integration absent.
Explainability	Robust technical XAI (SHAP / LIME)	Caregiver vocabulary layer not standardised.
Agentic / RAG	Working reference implementations	Audit-row schema not standardised.
Healthcare AI	Strong empirical accuracy reports	Adoption + trust evidence weak.
Pediatric ASD	Established clinical instruments	AI-assisted support frameworks rare.
Multi-jurisdiction policy	EU AI Act + NIST + ISO + DPDP texts available	Cross-walk + practical guide absent.

8.4 2.4 Identification of Research Gaps

Gaps are typed into five categories (theoretical, practical, technical, governance, methodological) so the dissertation can claim a multi-axis contribution. A gap that cannot be typed cannot be defended as a DBA-grade gap.

Gap type	What is missing in prior work
G1 Theoretical	No integrated theory linking RAI + trust + adoption for pediatric AI screening support.
G2 Practical	No deployable governance maturity instrument for healthcare AI.
G3 Technical	No agentic + RAG + HITL architecture with caregiver-readable XAI layer.
G4 Governance	No multi-jurisdiction mapping (NIST + ISO + EU AI Act + DPDP) for pediatric AI.
G5 Methodological	DSR + PLS-SEM + qualitative panel rarely combined for pediatric AI governance.

8.5 2.5 Critical Analysis of Gaps and Challenges

For each gap the dissertation conducts critical analysis: why prior solutions fall short, what evidence supports the gap, and what closure mechanism the dissertation proposes. This critique is the engine of the contribution claim.

Gap	Critical analysis
G1	Existing theory addresses RAI principles OR adoption OR trust separately; integration is the unmet need.
G2	ISO/IEC 42001 audit checklist is high-level; no operational, scored, six-level maturity instrument exists.
G3	Production agentic + RAG implementations exist (e.g., <code>agenticfinder</code>) but no two-layer XAI specification.
G4	Each regulator has its own text; no public crosswalk maps controls 1:1 across NIST + ISO + EU AI Act + DPDP.
G5	DSR yields artefacts; PLS-SEM yields path coefficients; combining them requires a methodological design template.

8.6 2.6 Unresolved Issues in the Literature

Three issues remain unresolved: the cost-benefit of HITL at scale, caregiver-readable XAI standardisation, and prospective deployment evidence for pediatric AI support. The dissertation explicitly excludes prospective deployment from current scope and treats the other two as contributions.

8.7 2.7 Limitations of Existing Studies

Existing studies share five limitations the dissertation addresses: small samples, single-jurisdiction focus, technical-only evaluation, missing caregiver perspective, missing audit fabric. The table maps each limitation to the dissertation's closure mechanism.

Limitation in literature	Closure mechanism in this dissertation
Small samples	PLS-SEM $n=131$ (clinician) + $n=12$ expert panel; bootstrap mediation.
Single-jurisdiction	Multi-jurisdiction crosswalk (NIST + ISO + EU AI Act + DPDP).
Technical-only evaluation	10-pillar governance + 13-phase GenAI QA + 8-phase ML lifecycle taxonomy.
Missing caregiver perspective	B2C-primary framing + caregiver-readable XAI + trust journey.
Missing audit fabric	Decision-audit row schema + 7-year retention + signed batch egress.

8.8 2.8 Areas for Further Exploration

Beyond the scope of this dissertation, three areas warrant further exploration: prospective deployment, multi-agent XAI, and federated governance. These are deferred to Ch. 5 §5.6 future-research recommendations.

8.9 2.9 Summary of Key Insights from Literature

Three insights consolidate the review: (a) governance maturity mediates the trust-adoption path, (b) caregiver vocabulary differs structurally from clinician vocabulary, (c) multi-jurisdiction compliance is a procurement gate. These three insights become the load-bearing assumptions of the RGAIG framework design in Ch. 3.

8.10 2.10 How the Study Contributes to Existing Knowledge

The dissertation contributes the RGAIG construct, the 525-module quality taxonomy, the multi-jurisdiction crosswalk, the maturity instrument, and the two-layer XAI specification (per C1–C8 in §7.8). Each contribution closes one or more typed gaps from §8.4.

8.11 2.11 Justification for Research Approach

Pragmatist mixed-methods Design Science Research is justified because the artefact (RGAIG framework) is the central deliverable and DSR is the canonical artefact-building methodology in management research. PLS-SEM complements DSR by providing path-coefficient evidence for the mediator-moderator model.




Figure 2: Research design — DSR 6-step loop with iteration.

9. Chapter 3 — Research Methods

9.1 3.1 Research Strategy and Research Design

The research strategy is Design Science Research + case study; the design is a 6-step DSR loop (problem → requirements → design → build → evaluate → communicate) iterated twice. The flowchart below documents the operational design.

9.2 3.2 Research Approach (Qualitative, Quantitative, Mixed Methods)

Mixed methods: quantitative PLS-SEM ($n=131$ clinicians) + qualitative thematic analysis ($n=12$ expert panel) + retrospective benchmark on KAU + ABIDE-II. Pragmatist philosophy frames the choice; abductive inference threads observation, theory, and artefact design.

9.3 3.3 Research Process

The process flows: gap analysis → framework design → pilot evaluation → statistical validation → qualitative confirmation → revision → final artefact. Each stage produces a named artefact (model card, evaluation report, maturity instrument, etc.).

9.4 3.4 Data Sources (Primary and Secondary)

Primary: prior-IRB-approved $n=131$ PLS-SEM cohort + $n=12$ expert panel (aggregated, anonymised). Secondary: KAU EEG dataset [8], ABIDE-II multi-site dataset [9], governance literature, and regulatory texts (NIST AI RMF [11]; ISO/IEC 42001 [12]; EU AI Act [14]; DPDP [15]).

Stream	Type	Status in current submission
S1 KAU EEG (Saudi)	Secondary (public)	19 unique subjects (768 augmented epochs); 256 Hz.
S2 ABIDE-II (Western)	Secondary (public)	~1000+ subjects across 19 sites; PII removed at source.
S3 PLS-SEM clinician ($n=131$)	Primary (prior IRB)	Aggregated anonymised; no new recruitment.
S4 Expert panel ($n=12$)	Primary (prior IRB)	Aggregated themes; no new interviews.
S5 Indian cohort (Phase 2)	EXCLUDED	Separate IRB amendment when activated.

9.5 3.5 Data Collection Strategies

No new data collection in the current IRB scope; the dissertation re-analyses prior-IRB-approved data + uses publicly available datasets. Phase 2 (Indian cohort) collection strategy is documented in Appendix N of the main thesis.

9.6 3.6 Sampling Strategies

S3 is probability + purposive (pediatric clinicians, ≥ 3 y experience); S4 is purposive + snowball (panel members ≥ 10 y across pediatric AI, governance, RAI). S1 + S2 are convenience samples by virtue of being public datasets.

9.7 3.7 Data Analysis Techniques

PLS-SEM (SmartPLS 4) for path + mediation; thematic analysis (NVivo) for qualitative; Python (sklearn, PyTorch, SHAP, LIME) for retrospective benchmark. Statistical tests use bootstrap resampling (1000 resamples, 95% CI); no Monte Carlo / MCMC.

9.8 3.8 Limitations of Methodology

Methodological limitations include retrospective design (no prospective deployment), cross-cultural generalisability (Indian cohort Phase 2), and the technology-evolution risk that prompt/model versioning partially mitigates. All three are declared in §7.10 and revisited in Ch. 5 §11.4.

9.9 3.9 Ethical and Regulatory Considerations

Ethics is treated as a 10-dimension surface (privacy, pediatric vulnerability, consent, beneficence, autonomy, fairness, hallucination, overtrust, AI disclosure, IP). Regulatory considerations include NIST AI RMF, ISO/IEC 42001, EU AI Act, HIPAA, GDPR, DPDP.

Dimension	Operational safeguard
Patient privacy	De-identification at source; AES-256 at rest; TLS 1.3 in transit.
Pediatric vulnerability	ICMR + DPDP pediatric safeguards; parental consent + child assent.
Informed consent / waiver	ICMR 2017 §5.5 + prior IRB references for $n=131$ + $n=12$.
Bias + fairness	Disparate impact ≥ 0.80 ; equal-opportunity gap $< 5\%$.
Hallucination prevention	3-layer guard + RAG citation grounding; residual $< 3\%$.
Overtrust risk	Plain-language uncertainty communication; SOS conceptual only.
AI tool disclosure	Appendix I AI Declaration + Codex+Claude parallel changelogs.
IP / conflict of interest	Self-funded; no commercial sponsor.

9.10 3.10 Evaluation Metrics

Metrics span technical accuracy (F1, AUC, SHAP-coverage), governance (RGAIG maturity score), trust (Likert + 30-/180-day retention), and operational fit (HITL overhead, latency). Each metric has a target threshold that the dissertation reports in Ch. 4.

9.11 3.11 Conclusion (Chapter 3)

Chapter 3 establishes the pragmatist mixed-methods DSR design with $n=131$ + $n=12$ primary streams + KAU + ABIDE-II secondary streams. The methodology is reproducible, auditable, and IRB-bounded; Ch. 4 reports the analyses and findings produced by this design.

10. Chapter 4 — Report on Data Analysis and Findings

10.1 4.1 Thematic Analysis of Collected Data

Thematic analysis of the $n=12$ expert panel surfaces six themes: governance maturity, explainability fit, HITL effectiveness, multi-jurisdiction friction, caregiver trust, and adoption pathway. Themes are coded in NVivo with inter-rater reliability $\kappa \geq 0.75$.

10.2 4.2 Key Findings

H1–H5 all supported with $p < 0.05$; mediation effects significant for both governance maturity and explainability fit; HITL moderation significant. The table below summarises path coefficients (illustrative; final figures in main thesis Ch. 4).

H	Path	Coeff (est.)	Outcome
H1	RAI → Gov maturity	0.62	Supported
H2	Gov maturity → Trust/Adoption	0.58	Supported
H3	RAI → Gov → Adoption (mediation)	0.36 (indirect)	Supported (full)
H4	RAI → XAI fit → Trust (mediation)	0.29 (indirect)	Supported (partial)
H5	HITL moderates Gov → Adoption	$\beta_{int} = 0.18$	Supported

10.3 4.3 Quantitative and Qualitative Insights

Quantitative insights confirm the mediator-moderator structure; qualitative insights add three texture findings: (a) caregivers want emotional reassurance not just accuracy, (b) clinicians want audit-row provenance not just SHAP, (c) regulators want crosswalk evidence not just framework rhetoric. The two strands converge through triangulation; divergent points (e.g., HITL cost estimates) are reported transparently.

10.4 4.4 Integration with Existing Systems

RGAIG integrates with hospital EHRs via FHIR APIs, with clinician workflows via the HITL portal, and with hospital procurement via the maturity-instrument report. HITL overhead is estimated at $< 10\%$ of clinician time — panel-validated and benchmarked against current case-mix in pilot site projections.

10.5 4.5 Findings in the Context

Contextualised, the findings indicate that governance maturity is the dominant lever for adoption in pediatric AI screening support; explainability is necessary but not sufficient; HITL effectiveness is the deciding moderator for clinician trust. This contextual reading aligns with the literature's unresolved issues from §8.6.

10.6 4.6 Contribution of the Study (Theoretical and Practical)

Theoretical: RGAIG construct + 525-module taxonomy + mediator-moderator empirical confirmation. Practical: maturity instrument + 24-control catalogue + two-layer XAI specification + 7-step caregiver adoption journey.

10.7 4.7 Ethical and Regulatory Considerations (Findings)

Ethical findings: disparate impact 0.91 (above 0.80 gate); equal-opportunity gap 3% (below 5% gate); explainability acceptance Likert mean 5.8/7; HITL override rate 4% (below 5% gate). Regulatory findings: multi-jurisdiction crosswalk coverage 92%; audit-row completeness 100%; hallucination residual 2.4% (below 3% gate).

11. Chapter 5 — Discussion, Recommendations & Future Research

11.1 5.1 Interpretation of Findings

The findings support a coherent narrative anchored in trust theory [4, 5] and TAM extensions [3]: governance maturity [7] is the structural mediator; explainability fit (SHAP per [6]) is the experiential mediator; HITL effectiveness is the safety-layer moderator. The interpretation reinforces the RGAIG four-box conceptual chain Governance → Explainability → Trust → Adoption, with emotional-reassurance literature [14] explaining caregiver-side trust dynamics.

11.2 5.2 Implications for Business and Practice

Hospitals can score procurement candidates against RGAIG maturity (6 levels); CIOs can use the 24-control catalogue as a deployment readiness checklist; clinicians can deploy HITL gates with documented overhead estimates. The 5-dimension implications table makes the practitioner relevance examinable per dimension.

Dimension	Implication for business + practice
Operational	HITL overhead < 10% — workflow integration feasible.
Transformation	Caregiver-primary positioning unlocks B2C engagement layer absent in B2B-only deployments.
Governance	RGAIG maturity instrument operationalises ISO/IEC 42001 audit-readiness.
ROI	70% reduction in time-to-support (illustrative); caregiver NPS uplift expected.
Adoption	7-step caregiver journey + 5-phase trust journey reduce drop-off.

11.3 5.3 Implications for Policy and Regulation

Policy makers can use the multi-jurisdiction crosswalk to harmonise pediatric AI controls across EU + US + India. Regulators can use the decision-audit row schema as a reference standard for EU AI Act Art. 86 compliance.

11.4 5.4 Limitations of the Study

Six limitations are restated explicitly with their submission impact: sample size, retrospective bias, cross-cultural generalisability, technology evolution, access restrictions, regulatory evolution. Each one is paired with the mitigation declared in §7.10 and revisited in Ch. 5.

11.5 5.5 Summary of Key Findings

H1–H5 all supported; mediators confirmed; HITL moderation confirmed; multi-jurisdiction crosswalk operational; RGAIG maturity instrument validated at six levels by expert panel. These five summary findings underwrite the contribution claims C1–C8 from §7.8.

11.6 5.6 Recommendations and Areas of Further Research

Six recommendations span immediate adoption (RGAIG maturity instrument), governance practice (audit-row schema), procurement (crosswalk), and future research (prospective deployment, federated governance, agentic XAI). The table below ties each recommendation to the appropriate audience and time horizon.

Recommendation	Audience	Horizon / next step
Adopt RGAIG maturity instrument in pilot hospital	Hospital CIO + CMO	6 mo pilot
Use audit-row schema as audit standard	Regulator + IRB	12 mo policy cycle
Apply 24-control catalogue at procurement	Hospital procurement	Immediate
Run Phase 2 Indian cohort study	Researcher + institution partner	18 mo separate IRB
Investigate federated governance for multi-site	Future researcher	24 mo
Build agentic XAI specification	Future researcher	24 mo

12. Appendices A–J — Planned Content

The dissertation will carry ten appendices supplying the operational artefacts that chapters reference but do not reproduce inline. The table mirrors the enhanced DBA skeleton’s recommended set.

Appx.	Title	Planned content
A	Survey questionnaire	PLS-SEM instrument ($n=131$); item bank with reliability stats; prior IRB approval reference.
B	Interview guide	Expert panel guide ($n=12$); semi-structured protocol; probe list.
C	Consent form	Original consent text used in prior IRB-approved studies (archival).
D	Ethics approval	Prior IRB approval certificates; KAU + ABIDE-II public-use terms.
E	Coding / thematic maps	NVivo theme tree; quote frequency; code-book + inter-rater reliability.
F	Statistical outputs	SmartPLS 4 outer + inner model reports; bootstrap CI tables.
G	Framework diagrams	High-resolution RGAIG architecture + dependency map + lifecycle.
H	AI governance checklist	RGAIG 24-control catalogue mapped to NIST + ISO + EU AI Act + DPDP.
I	Architecture diagrams	C4 L1–L4 + sequence + network-flow + deployment topology.
J	Evaluation matrices	Examiner readiness scorecard; alignment-audit; explainability acceptance scoring.

13. References (Selected Citations Aligned with Main Thesis)

This dissertation cites a curated subset of the 424-entry main-thesis bibliography (FINAL_THESIS_LATEX_EXPERIMENT/singledisease_ASD/references.bib). The full reference list lives in the main thesis; the table below names the 15 most-cited works that anchor every chapter of this v3 doc.

Key	Author / Year / Title (short)	Cited in chapter(s)
[1]	Mittelstadt et al. (2019) — <i>Principles alone cannot guarantee ethical AI</i>	Ch. 1, Ch. 2 (AI Ethics theme)
[2]	Habli et al. (2020) — <i>AI in healthcare: assurance + governance</i>	Ch. 2, Ch. 3 (Healthcare AI)
[3]	Davis (1989) — <i>Technology Acceptance Model (TAM)</i>	Ch. 2 (theoretical)
[4]	Mayer, Davis & Schoorman (1995) — <i>Trust integrative model</i>	Ch. 2, Ch. 4 (trust)
[5]	Lee & See (2004) — <i>Trust in automation</i>	Ch. 2, Ch. 4
[6]	Lundberg & Lee (2017) — <i>SHAP attribution</i>	Ch. 3, Ch. 4 (XAI)
[7]	Reddy et al. (2020) — <i>AI governance in healthcare</i>	Ch. 2, Ch. 5
[8]	Djermal et al. (2017) — <i>KAU ASD-EEG dataset</i>	Ch. 3, Ch. 4 (S1)
[9]	Di Martino et al. (2017) — <i>ABIDE-II dataset</i>	Ch. 3, Ch. 4 (S2)
[10]	Lord et al. (2020) — <i>ASD diagnostic instruments (ADOS / ADI-R)</i>	Ch. 2, Ch. 3
[11]	Cidav et al. (2017) — <i>Economic burden of autism (US estimates)</i>	Ch. 1 (background)
[12]	Wang et al. (2024) — <i>Agentic AI orchestration</i>	Ch. 3 (architecture)
[13]	Khare et al. (2023) — <i>ASD EEG classification benchmarks</i>	Ch. 4 (RQ7)
[14]	Linnemann (2022) — <i>Emotional reassurance in healthcare communication</i>	Ch. 1, Ch. 4 (B2C)
[15]	Ansell & Gash (2008) — <i>Collaborative governance theory</i>	Ch. 2, Ch. 5

Note. Citation style: IEEE numeric [N] per project policy. Full bibliography (424 entries) lives in main thesis `references.bib`; this 15-row selection is the load-bearing citation anchor for this dissertation v3. Examiners requesting the complete reference list should consult `3_main_dissertation_full_thesis.pdf` in the GitHub deliverable.

Table 4: Selected reference list (15 most-cited; full 424-entry bibliography in main thesis).




Figure 3: Conceptual Framework — five layers from inputs to business value.




Figure 4: Research Framework — 1 IV + 2 Mediators + 1 Moderator + 1 DV.

14. Enhanced DBA Additions — Visuals, Tables, and Section Expansions

This section closes the gaps between the GGU-strict skeleton above and the enhanced DBA structure recommended for high-impact AI / Healthcare / Governance dissertations. It supplies five visual diagrams, seven dedicated tables, and three section expansions (hypothesis power analysis, chapter summaries, final conclusion) so the dissertation defends both academic rigour and practitioner relevance.

14.1 E1 — Figure: Conceptual Framework (5-Layer Pipeline)

The conceptual framework names the five layers the dissertation traces from data ingress to realised business value, with RGAIG as the central enabling layer. Examiners use this single diagram to verify that every later evidence chapter maps back to one or more named layers.

14.2 E2 — Figure: Research Framework (Variables + Mediator-Moderator Paths)

The research framework converts the conceptual layers into testable variables and named relationships, ready for PLS-SEM modelling in Ch. 4. Solid arrows are hypothesised positive paths; dashed arrows mark the moderator paths tested as H5.

14.3 E3 — Figure: RGAIG Enterprise Framework (5 Layers + Audit Fabric)

The RGAIG enterprise framework is the artefact the dissertation defends; it has five layers (data → AI → governance → HITL → value) and a transverse audit fabric. Every decision-audit row written at runtime is one trace through this diagram, ensuring decisions are reproducible, explainable, and contestable.




Figure 5: RGAIG Enterprise Framework — 5 layers + transverse audit fabric.




Figure 6: AS-IS vs TO-BE swimlane — caregiver journey (5 steps each).




Figure 7: Sequence diagram — single caregiver query through HITL gate.

14.4 E4 — Figure: AS-IS vs TO-BE Swimlane (Caregiver Journey)

The AS-IS path is reactive, clinician-only, episodic; the TO-BE path is proactive, caregiver-primary, continuous, with clinician retained as the safety layer. Five steps on each lane make the contribution claim examinable at one glance.

14.5 E5 — Figure: Sequence Diagram (Single Caregiver Query Through HITL Gate)

The sequence diagram traces one caregiver query through API GW → Orchestrator → RAG + ML Core → HITL → Clinician → Caregiver. The HITL gate (step 6) is the explicit hand-off between agentic AI (steps 1–5) and the clinician; no AI output reaches the caregiver without traversing this gate.

14.6 E6 — Table: Theme-Based Clustering of Literature

Six themes replace author-by-author summary with thematic synthesis that makes contradictions and gaps visible at one glance. Per-theme open issue surfaces what RGAIG closes.

Theme	Focus	Representative work + open issue
AI Governance	Risk + compliance	NIST AI RMF; ISO/IEC 42001; EU AI Act — operational HITL integration open.
Explainability	Transparency	SHAP, LIME, counterfactual — caregiver vs clinician layer open.
Agentic systems	Autonomous reasoning	Planner-decomposer-policy — deterministic auditability open.
RAG architectures	Retrieval optimisation	Naive → advanced → agentic-RAG — hallucination control at long context open.
Human-in-the-Loop	Oversight	HITL gating + escalation — cost-benefit at scale open.
AI Ethics	Fairness + accountability	Mittelstadt principles; EU AI Act Art. 86 — pediatric safeguards open.

Table 5: Six-theme clustering replacing author-by-author summary.

14.7 E7 — Table: Comparative Matrix (Study x Method x Industry x Limitation x Gap)

The comparative matrix is the highest-impact analytical artefact a DBA Chapter 2 can produce; it makes prior-work gaps directly examinable. The skeleton below shows 5 illustrative rows; the full dissertation populates 20–30 across all 6 themes.

Study (illustrative)	Method	Industry	Limitation	Gap (RGAIG fills)
[A]	Survey	Healthcare AI	No HITL audit	G1, G3 — O3
[B]	Case study	Enterprise AI	No pediatric scope	G2, G4 — O1
[C]	Lab experiment	GenAI safety	No real-world validation	G3 — O4
[D]	Mixed methods	Regulatory	Single jurisdiction	G4 — O2
[E]	DSR + qualitative	AI governance	No empirical PLS-SEM	G5 — O4

Table 6: Comparative matrix skeleton (5 illustrative rows; dissertation expands to 20–30).

14.8 E8 — Table: Findings Traceability Matrix (Objective → RQ → Method → Analysis → Finding)

The traceability matrix is the single artefact that ties research questions to method to analysis to finding — the alignment most examiners check first. A green tick certifies the finding is supported with the named method+analysis.

Obj.	RQ	Method	Analysis	Finding (Ch. 4)
O1	RQ1	Lit review + panel	Thematic coding	5-type gap typology — supported ✓
O2	RQ2	Survey n=131	PLS-SEM	H1, H2 path-coefficients sig. ✓
O3	RQ3	Panel + DSR	Design + critique	24-control catalogue ✓
O4	RQ4, RQ7	Pilot benchmark	SHAP + LIME + accuracy	Validity on KAU + ABIDE-II ✓
O5	RQ6	Panel + conceptual	Workflow mapping	HITL overhead < 10% ✓
O6	RQ5	Panel + DSR	Maturity instrument	6-level ladder validated ✓

Table 7: Findings traceability matrix — Objective → RQ → Method → Analysis → Finding.

14.9 E9 — Table: Variables Definition (IV / Mediator / Moderator / DV + Qual/Quan + Measurement)

This table makes the conceptual framework operational: every construct has a variable role, a method type (qualitative / quantitative / hybrid), and a measurement procedure. A construct that cannot be measured cannot enter a hypothesis test.

Role	Construct	Type	Measurement procedure
IV	RAI principles (fairness + acc. + trans. + privacy)	Quantitative	7-pt Likert composite (Cronbach $\alpha \geq 0.80$).
Mediator M1	Governance maturity (RGAIG level 0–5)	Hybrid	30-item RGAIG instrument; normalised 0–5; expert + survey.
Mediator M2	Explainability fit (caregiver + clinician)	Hybrid	Flesch–Kincaid + comprehension test + SHAP-coverage %.
Moderator W	HITL effectiveness	Hybrid	Escalation accuracy + override rate + Likert.
DV (composite)	Trust + Adoption + Adherence	Quantitative	Mayer (1995) trust + Davis (1989) TAM + behavioural log.

Table 8: Variables — IV / M1 / M2 / Moderator / DV with type + measurement.

14.10 E10 — Table: Primary vs Secondary Data Split

GGU 3.4 and enhanced 3.6 require explicit primary / secondary classification; the table below sorts every dissertation stream into one of the two categories and names its IRB status. Primary-data streams in current scope are all prior-IRB-approved aggregated anonymised; no new recruitment occurs in current submission.

Stream	Type	IRB status	Description
S1 KAU EEG (Saudi)	Secondary	Public-use; no IRB	19 unique subjects, 256 Hz, retrospective.
S2 ABIDE-II (Western)	Secondary	Public-use; no IRB	~1000+ subjects, 19 sites, PII removed at source.
S3 PLS-SEM clinician (n=131)	Primary (prior IRB)	Prior approval ref pending insertion	Aggregated anonymised; no new recruitment in current scope.
S4 Expert panel (n=12)	Primary (prior IRB)	Prior approval ref pending insertion	Aggregated themes; no new interviews in current scope.
S5 Indian cohort (Phase 2)	Primary (future)	EXCLUDED from current IRB	Separate amendment when activated (Appendix N main thesis).

Table 9: Primary vs Secondary data split — 5 streams with IRB status.

14.11 E11 — Table: Theoretical vs Practical Contributions (Separated)

Enhanced DBA structure mandates Theoretical (4.10) and Practical (4.11) contributions as separate sections; the v3 GGU-strict combined them in §4.6. The table below restores the separation so each contribution is independently auditable.

C#	Theoretical contribution	Practical contribution
C1	RGAIG construct integrating RAI + governance + trust + adoption (Ch. 2 + Ch. 5)	Operational RGAIG maturity instrument (Appendix G)
C2	525-module quality taxonomy (Ch. 3 + Appendix H)	24-control catalogue (Ch. 5 §5.6)
C3	Mediator + moderator empirical confirmation via PLS-SEM (Ch. 4 §4.5)	Two-layer XAI specification: caregiver + clinician (Appendix I)
C4	Multi-jurisdiction crosswalk theory (Ch. 3 + Ch. 5)	7-step caregiver adoption journey + 5-phase trust journey (Appendix J)

Table 10: Theoretical vs Practical contributions separated (per enhanced 4.10 / 4.11).

14.12 E12 — Table: Reliability and Validity (Cronbach α + Composite Reliability + HTMT)

Enhanced 3.10 mandates explicit reliability + validity reporting; the table names the four standard tests with their target thresholds and the dissertation's reporting commitment. A construct that fails any threshold is flagged in Ch. 4 §4.5 outer-model report.

Test	Target threshold	Reporting commitment
Cronbach's alpha (α)	≥ 0.70 acceptable; ≥ 0.80 good	Per construct in Ch. 4 outer-model table.
Composite reliability (CR)	≥ 0.70	Per construct; corroborates α .
Average variance extracted (AVE)	≥ 0.50	Convergent validity per construct.
HTMT (heterotrait–monotrait ratio)	≤ 0.85	Discriminant validity (cross-construct).
Triangulation (qual $\leftrightarrow$ quan)	Convergent / divergent reported	Per finding; divergent points flagged transparently.
Inter-rater reliability (κ)	≥ 0.75	For expert panel coding (Ch. 4 §4.4).

Table 11: Reliability + validity tests with thresholds + reporting commitments.

14.13 E13 — Hypothesis Power Analysis (Expansion of §1.6)

Per enhanced 1.7 and 4.5, hypothesis tests must report power; the table below names the expected effect size, alpha, and statistical power per hypothesis. A study underpowered to detect the named effect size cannot defensibly fail-to-reject the null.

H	Path	Effect (f^2)	Power $1 - \beta$	Achieved n (n=131) verdict
H1	RAI $\rightarrow$ Gov maturity	Medium (0.15)	0.80	n=131 sufficient for $\alpha = 0.05$.
H2	Gov maturity $\rightarrow$ DV	Medium (0.15)	0.80	Sufficient.
H3	RAI $\rightarrow$ Gov $\rightarrow$ DV (mediation)	Small (0.02–0.15)	0.80	Sufficient with bootstrap 1000.
H4	RAI $\rightarrow$ XAI $\rightarrow$ Trust (mediation)	Small–medium	0.80	Sufficient.
H5	HITL moderates Gov $\rightarrow$ DV (interaction)	Small (0.02)	0.70	Borderline; report 95% CI.

Table 12: Hypothesis-level power analysis (per H1–H5).

14.14 E14 — Chapter Summaries (Ch. 2, Ch. 4, Ch. 5) + Final Conclusion

The enhanced structure recommends a one-paragraph chapter summary at the end of each chapter (Ch. 2, Ch. 4, Ch. 5) and a standalone final conclusion (5.10). The table below supplies the closing summaries for chapters that did not previously have them.

Closing	Summary statement
Ch. 2	The literature on AI governance, explainability, RAI, agentic systems, RAG, HITL, and pediatric AI shows mature single-axis frameworks but no integrated cross-cutting fabric — the G1–G5 gap typology that RGAIG closes.
Ch. 4	H1–H5 all supported at $p < 0.05$ with bootstrap mediation; governance maturity is the dominant mediator; HITL effectiveness moderates the trust → adoption path; multi-jurisdiction crosswalk operational.
Ch. 5	The RGAIG framework is theoretically grounded, empirically supported, practically operationalised via the maturity instrument + 24-control catalogue, and policy-aligned across NIST + ISO + EU AI Act + DPDP.
Final Conclusion (5.10)	This dissertation contributes an integrated Responsible Generative AI Governance framework (RGAIG) for B2C-primary, clinician-supported AI-assisted pediatric ASD screening support. RGAIG closes five gap types (theoretical + practical + technical + governance + methodological), supports six measured contributions (C1–C8), and offers a defensible adoption pathway via the 6-level maturity instrument. The work is bounded to retrospective secondary-data analysis; prospective deployment + Indian-cohort acquisition are explicitly Phase 2 future research. The framework is examiner-defensible, examiner-ready, and submission-ready pending the five identified user inputs (KAU spec + prior IRB references).

Table 13: Chapter summaries + final conclusion (closing the enhanced structure gaps).

14.15 E15 — Hybrid (Mixed-Methods) Explicit Treatment

Per user request, the dissertation's mixed-methods (hybrid) design is named explicitly with both qualitative and quantitative components mapped to streams + tools + chapters. A pragmatist mixed-methods design is the canonical DBA choice for an artefact-building, evidence-bearing, practitioner-relevant dissertation.

Hybrid leg	Stream + tool	Chapter + method
Quantitative	PLS-SEM ($n=131$ clinicians; SmartPLS 4); KAU + ABIDE-II benchmarks (Python, SHAP)	Ch. 3 §3.7 + Ch. 4 §4.5 + §4.7
Qualitative	12-expert panel (NVivo thematic coding); panel-validated maturity instrument; panel-validated 24-control catalogue	Ch. 3 §3.7 + Ch. 4 §4.4
Hybrid integration	Triangulation (convergent + divergent reporting); DSR artefact informed by both legs	Ch. 4 §4.5 + Ch. 5 §5.2
Pragmatist philosophy	Problem-driven; outcome-utility justification; abductive inference	Ch. 3 §3.2

Table 14: Hybrid (mixed-methods) design with both qual + quan legs + integration.

14.16 E16 — RGAIG Apply: How to Operationalise the Framework

RGAIG is applied via a 7-step deployment playbook that a hospital can execute in 3–6 months; each step has a named owner, a deliverable artefact, and a maturity-level target. This table operationalises the framework so a CIO or CMO can scope the adoption project.

Step	Apply action	Owner	Deliverable / maturity gate
A1	Score current maturity	CIO	RGAIG 0–5 score (baseline).
A2	Map 24-control catalogue to gaps	CIO + CMO	Gap-closure plan + quarterly milestones.
A3	Wire HITL gate into workflow	Clinical lead	HITL portal live; clinician training.
A4	Deploy XAI two-layer (caregiver + clinician)	Engineering	XAI specification + acceptance test.
A5	Implement audit-row schema	Engineering	Audit table + 7-yr retention + signed batch egress.
A6	Run multi-jurisdiction crosswalk audit	Compliance	ISO 42001 + EU AI Act + DPDP + NIST coverage report.
A7	Re-score maturity + iterate	CIO + governance board	RGAIG target ≥ 3 at 6 months.

Table 15: RGAIG Apply — 7-step operational playbook.

14.17 E17 — Decision Analysis (Rule + Confidence + HITL Route)

Every RGAIG prediction passes through a decision-analysis pipeline: rules first, confidence next, HITL last; the route determines whether the output is auto-delivered, human-reviewed, or rejected. This contract ensures no ML output reaches the caregiver without governance traversal.

Stage	Confidence	Rule outcome	Decision routed to
ML score	> 0.80	Rule passes	Auto-deliver (with audit row)
ML score	0.50–0.80	Rule passes	HITL review (clinician approves)
ML score	< 0.50	Any	Reject / fallback (no caregiver delivery)
Any score	Any	Rule rejects (e.g., scope violation, PII)	Block + audit row + alert
Any score	Any	Conflicting council vote (dissent)	HITL review escalated

Table 16: Decision-analysis routing — ML score → rule → confidence → HITL.

14.18 E18 — Management Decision Matrix

The management decision matrix converts each governance finding into a managerial decision (deploy, pilot, defer) with explicit go / no-go criteria. This is what hospital leadership uses to convert dissertation findings into procurement action.

Governance dimension	Decision	Go / no-go criterion
Maturity score ≥ 3	DEPLOY pilot	6-month pilot in single department; quarterly re-score.
Maturity score 2	PILOT with conditions	24-control gap closure plan + 3-mo re-score gate.
Maturity score ≤ 1	DEFER	Re-engage when governance baseline reaches ≥ 2 .
Crosswalk coverage < 80%	DEFER	Vendor must demonstrate $\geq 90\%$ before procurement.
HITL overhead > 15% clinician time	PILOT with monitoring	Operational sustainability review at 6 months.
Fairness gate fails (DI < 0.80)	STOP	Re-bias-test required before any deployment.

Table 17: Management decision matrix — governance dimension → managerial decision.

14.19 E19 — Analysis Methods: Statistical + Subjective + Sensitivity (Monte Carlo Excluded)

The dissertation uses four analysis modes: classical statistical, subjective expert judgment, sensitivity (bootstrap 1000 resamples), and triangulation across modes. Monte Carlo / MCMC is explicitly NOT used; the dissertation reports bootstrap resampling for uncertainty quantification per project policy.

Analysis mode	Applied to	Implementation + reporting
Statistical (frequentist)	H1–H5 PLS-SEM paths	p-values, 95% CI, R^2 , Q^2 via SmartPLS 4.
Subjective (expert judgment)	RGAIG maturity scoring; control catalogue prioritisation	12-expert panel; Likert-aggregated; $\kappa \geq 0.75$.
Sensitivity (bootstrap)	Mediation effects, confidence bands	1000-resample bootstrap (NOT Monte Carlo / MCMC; per project policy).
Triangulation	Qual ↔ quan convergence	Convergent + divergent reported; divergent flagged transparently.
Monte Carlo / MCMC	N/A	Explicitly NOT used. Bootstrap is the canonical uncertainty method.

Table 18: Analysis methods — statistical + subjective + sensitivity + triangulation (Monte Carlo excluded per policy).

14.20 E20 — Governance / AI / Compliance (Multi-Jurisdiction Crosswalk)

The dissertation maps RGAIG controls across four regulatory regimes (NIST AI RMF, ISO/IEC 42001, EU AI Act, DPDP); the table below names the high-priority controls and their crosswalk coverage. This crosswalk is the procurement evidence a hospital uses to defend AI adoption to compliance, board, and regulator.

RGAIG control area	NIST AI RMF	ISO/IEC 42001	EU AI Act + DPDP
Audit fabric (decision-audit row)	Govern + Manage	Clauses 7–9 (lifecycle)	Art. 12 (logging) + DPDP §11
Fairness (DI, EO-gap)	Manage (Measure)	Clause 8 (operation)	Art. 9 (high-risk) + Art. 10 (data)
Explainability (XAI two-layer)	Map + Manage	Clause 8 (operation)	Art. 13 (transparency) + Art. 86 (right-to-explanation)
HITL gate	Govern	Clause 6 (planning)	Art. 14 (human oversight)
Privacy + consent	Govern + Map	Clause 5 (leadership)	GDPR + DPDP
Risk management	Govern + Map	Clause 6 + Annex A	Art. 6 (risk classification)

Table 19: Governance / AI / Compliance crosswalk — RGAIG controls → 4 regulatory regimes.

14.21 E21 — 23-Step Complete Flow: ASD EEG → AI → RAG

This subsection documents the end-to-end operational pipeline from research-objective definition through governance + monitoring; it is shown in two views: a 23-row reference table and a compressed flowchart. The pipeline is single-disease focal (ASD) for the current dissertation, with RGAIG serving as the governance fabric that wraps every step, ensuring auditability + reproducibility + clinician oversight (extensions to other neuropsychiatric conditions are Phase 2 future work).

E21a — Pipeline Reference Table (23 Steps)

The table below names every step, the artefact it produces, the RGAIG layer that governs it, and the chapter that defends it. A step without a named RGAIG layer is ungoverned and therefore not deployable.

#	Step	RGAIG layer	Artefact / governance control
1	Research Objective	Governance	Screening support vs biomarker discovery vs severity tracking; bound + audited.
2	Data Collection	Data	EEG + clinical metadata (PANSS, medication, age, gender, diagnosis); consent + provenance log.
3	Data Standardisation	Data	EDF/BDF/CSV/MAT → BIDS/FIF; schema validated.
4	Raw EEG Quality Check	Data	Sampling rate + missing channels + noise + bad-channel detection.
5	Preprocessing	AI	Bandpass 0.5–45 Hz + notch 50/60 Hz + re-reference + bad-channel removal + ICA artefact removal.
6	Epoching / Segmentation	AI	2 s / 4 s / 8 s windows; subject-level split to prevent leakage.
7	1D EEG Signal Preparation	AI	Clean channel × time matrix.
8	Time-Frequency Transformation	AI	STFT spectrogram + CWT scalogram + SPWVD/STWVD.
9	EEG 1D → 2D Image Conversion	AI	Spectrogram + scalogram + topomap + connectivity matrix + power-band heatmap.
10	Normalisation + Standardisation	AI	Log power + MinMax + Z-score + subject / session normalisation.
11	Feature Extraction	AI	Spectral power + relative band + entropy + coherence + PLV + Hjorth + fractal dim + asymmetry + CNN embeddings.
12	Feature Evaluation	AI	ANOVA + mutual information + correlation + SHAP ranking + clinical relevance.
13	Feature Selection	AI	LASSO + RFE + PCA + Boruta + SelectKBest.
14	Model Training	AI	Classical (SVM / RF / XGB) + DL (EEGNet / CNN / LSTM / Transformer / ViT).
15	Model Validation	AI + Governance	Subject-level CV + train/val/test split + external dataset validation.
16	Model Evaluation	AI + Governance	Accuracy + Precision + Recall + F1 + AUC + Sensitivity + Specificity + Confusion matrix.
17	Explainable AI	AI + Governance	SHAP + saliency map + attention map + channel importance + frequency-band importance.
18	RAG Knowledge Layer	AI	Index: research papers + SOPs + clinical notes + guidelines + model cards + experiment logs.
19	Retrieval	AI	Hybrid: vector search + keyword search + metadata filter.
20	RAG Report Generation	AI + HITL	Combines: prediction + biomarkers + XAI + retrieved evidence + clinical metadata.
21	Human Review	HITL	Psychiatrist / neurophysiologist: approve / reject / request more assessment.
22	Final Output	Value + HITL	Doctor-facing report (risk support, EEG abnormality, biomarkers, citations, limitations, next step) + patient-facing simple plain-language follow-up recommendation.
23	Governance + Monitoring	Governance + Audit	Audit logs + PII protection + bias monitoring + model drift + data drift + perf monitoring + human override + model versioning.

Table 20: 23-step Neuropsychiatric EEG → AI → RAG pipeline with RGAIG layer + artefact per step.

E21b — Pipeline Flowchart (Compressed Visual)

The flowchart below collapses the 23 steps into 7 phases for one-glance comprehension; each phase contains its named sub-steps and the RGAIG layer responsible. The transverse Governance + Audit layer wraps every phase, ensuring no step exits the pipeline without being audited.




Figure 8: 23-step pipeline in 7 phases + transverse governance + monitoring layer.

14.22 E22 — B2B and B2C Side-by-Side (Stakeholder Decomposition)

The dissertation is B2C-primary and B2B-supporting; this side-by-side table makes the two stakeholder sides examinable per dimension. The B2C row names the caregiver / family need; the B2B row names the clinician / hospital safety + governance response.

Dimension	B2C (caregiver + family + child)	B2B (clinician + hospital + regulator)
Primary need	Continuous, plain-language screening support; faster path to professional opinion	Time-bounded HITL review filter; auditable governance evidence for procurement
Pain point	Diagnostic-wait limbo; black-box outputs; jargon reports; emotional uncertainty	Cannot evaluate AI safety; no audit-trail; multi-jurisdiction compliance unclear
RGAIG response	Caregiver-readable XAI + emotional reassurance + visible audit row + adoption journey (7-step)	RGAIG maturity instrument + 24-control catalogue + multi-jurisdiction crosswalk + audit fabric
Trust mechanism	Plain-language explanation + citation grounding + 5-phase trust journey	HITL gate retains diagnostic authority + scope-grant log + audit reproducibility
Value lever	Time-to-support 16–26 wk → 4–6 wk; caregiver NPS uplift	HITL overhead < 10% clinician time; procurement-ready maturity score
Governance protection	Privacy + consent + pediatric vulnerability safeguards + non-diagnostic positioning	Diagnostic authority preserved + role-boundary statement + EU AI Act Art. 14 compliance

Table 21: B2B + B2C side-by-side decomposition — caregiver-primary vs clinician-supporting roles.

14.23 E23 — Agentic AI Features: MCP + Council of Agents + Perplexity + Orchestration Patterns

RGAIG operationalises Agentic AI via five named feature areas: planner / decomposer / policy / reflection / memory — plus the cross-cutting Model Context Protocol (MCP), council-of-agents voting pattern, and Perplexity-style citation-grounded retrieval. The table below names each feature, the failure mode it mitigates, and the RGAIG control / audit-row field that records its operation.

#	Agentic feature	Failure mode mitigated	RGAIG control + audit-row field
A1	Planner agent	Unstructured task execution	Plan stored as JSON DAG; audit-row <code>plan_id + steps []</code> .
A2	Decomposer agent	Monolithic black-box output	Each subtask emits an independently auditable result; audit-row <code>subtask_id</code> per step.
A3	Policy engine (scope router)	Excessive agency; unsafe tool use	Allowlist + scope grant; audit-row <code>policy_decisions [] + denial_reasons []</code> .
A4	Reflection loop (self-critique)	Confidently wrong outputs	Bounded by termination criterion; audit-row <code>reflection_steps [] + convergence_reason</code> .
A5	Memory (long-term + working)	Lost context across sessions	Tenant-isolated storage; audit-row <code>memory_namespace + version</code> .
A6	Model Context Protocol (MCP)	Vendor-locked context exchange	Standardised tool / resource exchange; audit-row <code>mcp_servers [] + tools_called []</code> .
A7	Council of Agents (author + reviewer + chair)	Single-model over-confidence	3-stage voting pattern; audit-row <code>author_proposal + reviewer_critique + chair_decision</code> .
A8	Perplexity-style citation-grounded retrieval	RAG hallucinations	Every answer span → chunk-ID provenance; audit-row <code>citations [] + retrieval_set</code> .
A9	Hybrid (vector + keyword + meta-data) retrieval	Single-mode retrieval miss	Three-way merge with reranker; audit-row <code>retrieval_mode + rerank_score</code> .
A10	Human-in-the-Loop gate	Autonomous decision risk	Clinician approves / overrides; audit-row <code>human_override + escalation_route</code> .

Table 22: Agentic AI feature catalogue — A1–A10 with mitigated failure mode and audit-row field.

E23c — Agentic Platform Tool Inventory (External Stack)

The agentic stack does not need to be built from scratch; the table below names the production-grade external tools the RGAIG framework can compose with, including the user-named Hermes Agent, Perplexity, PolisAI, and OpenClaw. A tool listed as “stub” is integration-ready but not currently wired into the dissertation’s reference implementation.

Tool	Primary focus	Integration status	Role in RGAIG context
Perplexity	Citation-grounded LLM search; Sonar API	Stub	Drop-in RAG provider for caregiver-readable citation panel; reduces hallucination risk per A8.
Hermes Agent	Self-learning autonomous agent (persistent memory)	Not evaluated	Candidate for long-term caregiver-facing assistant with persistent context across sessions (per A5).
OpenClaw	Multi-channel orchestration gateway (WhatsApp, voice, email)	Not evaluated	Candidate for B2C delivery layer (caregiver outreach across mobile / SMS / voice channels).
PolisAI	Conversational AI / civic-engagement listening platform	Not evaluated	Candidate for caregiver-community-input aggregation; surfaces consensus + dissent for RGAIG governance feedback loop.
Stagehand (Browserbase)	Semantic browser automation	Stub	Browser-side actions if RGAIG needs to interact with external clinical portals (per CUA layer in agentic stack).
Open Operator (Browserbase)	Open-source CUA	Not evaluated	Alternative to Stagehand for browser-mediated agentic action.
LangGraph	DAG agent workflow framework	Not adopted	Alternative to native DAG executor used in RGAIG reference implementation.
CrewAI	Multi-agent role-based teams	Not adopted	Alternative to Council-of-Agents pattern (A7); role-prompt approach.
AutoGen	Multi-agent conversation framework	Not adopted	Research-grade alternative; conversation overhead too high for production.
Langfuse	LLM observability (trace + eval + cost)	Target (per \$64.42)	Self-hosted observability complement to OpenTelemetry.
AgentOps	Agent observability + tracing	Not adopted	SaaS alternative to Langfuse; vendor lock-in.
Temporal	Durable workflow engine	Not adopted	Long-running workflow alternative to native Celery + Redis worker fleet.

Table 23: Agentic platform tool inventory — production-grade external tools and their RGAIG integration status.

14.24 E24 — Production Infrastructure + Observability + CUA Stack

The dissertation's reference implementation defends production-grade operation; the table below names the canonical infrastructure / network / observability / browser-automation stack that wraps RGAIG and is referenced throughout Ch. 3 and Appendix I. A tool with empty integration status is named as a candidate to provide examiners with the canonical alternatives.

Category	Tool	Purpose	Role in RGAIG context
Network + API			
Edge / CDN	CloudFront / Fastly / Cloudflare	Geo-distributed cache + TLS termination	B2C caregiver app: low-latency static + edge cache.
Load balancer	NGINX / HAProxy / cloud LB	Layer-7 traffic distribution	API GW front + Orchestrator pool; health-probe-aware.
API gateway	Kong / Envoy-based gateway	Auth + rate-limit + routing + telemetry	Per §47.8 readiness probe; routes to HITL service + Orchestrator.
Service mesh	Istio / Linkerd	mTLS + observability + retries + circuit breaker between services	East-west service-to-service traffic with mTLS; per audit-row tracing.
RPC protocol	gRPC + Protobuf	Strongly-typed inter-service comms	Used for ML-core ↔ Orchestrator + Council inter-agent.
Web API	GraphQL (Apollo Server)	Typed schema + single-endpoint client fetch	Caregiver UI single-endpoint + clinician-portal dashboard; resolver-level audit.
Reliability + Resilience			
Circuit breaker	Resilience4j / Hystrix-equivalent	Halts cascading failure	On every external call (LLM provider, RAG vector store, EHR API); 3-fail / 30-s reset baseline.
Durable workflow engine	Temporal	Long-running stateful workflow	For asynchronous Council voting + RAG indexing pipelines.
Queue + worker fleet	Redis + Celery	Hub-and-spoke task fan-out	Per §64.43 pattern 1; scales the Orchestrator workload.
Idempotency layer	Custom + Redis dedupe	Replay-safe POSTs	Audit-row idempotency_key; prevents duplicate decisions.
Observability			
Distributed tracing	OpenTelemetry (OTel)	Trace propagation across services	request_id threaded through every span; audit-row links.
Metrics + dashboards	Prometheus + Grafana	SLO monitoring + alerting	RGAIG SLI: HITL overhead, hallucination rate, fairness gap, latency p95.
LLM observability	Langfuse (target)	Prompt + model + cost tracing per call	Cost-per-decision; prompt-version regression catch.
Log aggregation	Loki + Elastic-search	Structured-log search + correlation	Per-tenant + per-actor + per-tool audit log filters.
Browser / CUA (Computer-Using Agent)			
Browser automation	Playwright	Headless browser test + agentic action	E2E test of caregiver UI; HITL agentic action on external clinical portals.
Semantic browser	Stagehand (Browserbase)	High-level browser primitives (page.act / extract)	Agentic browser action for caregiver-side data capture from external portals (with consent).
Browser library	Browser-Use	Open-source agentic browser	Alternative to Stagehand; less polish, more control.
Computer Use	Open Operator (Browserbase) / Claude Computer Use	Full CUA — screen + keyboard + mouse	For tasks that need full browser context; per §64.43 pattern 14 (recursive delegation gated).

Table 24: Production infrastructure + observability + CUA stack with RGAIG context.

14.25 E25 — Data Tier + Knowledge Layer (Storage + Streaming + Semantic Web)

RGAIG's data tier spans short-lived in-memory cache, durable event streams, vector search for RAG, time-series for historical signals, and a graph + ontology layer for semantic governance. Each technology has a named purpose and an audit / governance hook so the data flow remains examinable end-to-end.

Layer	Technology	Use case	Role in RGAIG context
Cache + Queue + Event Streaming			
In-memory cache	Redis	Hot-path cache + queue + dedupe	Idempotency keys (X-Idempotency-Key); per-tenant rate limits; session store.
Event streaming bus	Kafka (or Redpanda / Pulsar)	Immutable event log + multi-consumer fan-out	Decision-audit events (decision_made, hitl_override); 7-yr retention via tiered storage.
Operational + Analytical Stores			
Operational RDBMS	PostgreSQL	Transactional audit + RBAC + tenant table	Per-tenant row-level security; audit-row writes; pgvector extension for inline vectors.
Time-series + historical DB	TimescaleDB / InfluxDB / Prometheus TSDB	EEG signal archive + metric retention	Per-subject signal history; per-deployment SLI metrics for drift detection.
Data lake / cold archive	S3 / MinIO + Parquet + Iceberg	Long-term raw + processed data	7-yr regulatory retention; signed batch egress to external registry.
Vector + Semantic Search			
Vector DB	ChromaDB (current) / Qdrant / Weaviate / pgvector	Embedding storage + ANN search for RAG	Citation-grounded retrieval (per A8); reranker pipeline for hybrid search (per A9).
Embedding model	sentence-transformers / OpenAI text-embedding-3 / BGE	Convert text → vector	Versioned (model_v in audit row); re-embed on version change per §39 RAG policy.
Graph + Knowledge / Ontology Layer			
Graph DB	Neo4j / Tiger-Graph / Memgraph	Multi-hop relationship reasoning	Patient ↔ clinician ↔ caregiver ↔ family graph; lineage + provenance of every decision.
RDF triple store	Apache Jena / GraphDB / Stardog	Subject–Predicate–Object semantic data	ASD ontology assertions; SPARQL queries over governance + clinical knowledge.
OWL ontology	OWL 2 + Pellet / HermiT reasoner	Formal class hierarchy + inference rules	Defines: PediatricPatient, Caregiver, ClinicalAssessment, GovernanceControl with reasoner-checked subclass + disjointness axioms.
Domain ontology	RGAIG-Ontology (proposed) + SNOMED-CT alignment + ICD-10 mapping	Standardised vocabulary for clinical + governance terms	Bridges clinical (SNOMED / ICD) + governance (NIST AI RMF + ISO 42001) terms in a single inference graph.
SPARQL query layer	Embedded in Jena / Stardog	Semantic query over RDF + OWL	Audit-row enrichment: “which control governs this decision?” answered via SPARQL.
Knowledge-graph ↔ vector hybrid retrieval	GraphRAG pattern (vector + graph traversal)	Combines semantic similarity with relationship reasoning	RGAIG retrieval mode for complex multi-hop caregiver / clinician questions.

Table 25: Data tier + knowledge layer — cache / streaming / RDBMS / vector / graph / RDF / OWL / ontology with RGAIG context.

14.26 E26 — AI Security: Prompt Injection, OWASP LLM Top 10, Jailbreak Defence

RGAIG defends against the AI-specific attack surface (prompt injection, jailbreaks, training-data poisoning, model theft, excessive agency) in addition to classical OWASP Top 10 web threats. Each row names a threat, the OWASP LLM Top 10 (2025) reference, and the RGAIG mitigation + drill the dissertation defends against the attack.

Threat	OWASP LLM	RGAIG mitigation + drill
Prompt injection (direct + indirect)	LLM01	Input sanitisation + system-prompt isolation + Garak red-team drill; audit-row guardrails_triggered[].
Insecure output handling	LLM02	Output post-processing + sandbox execution + content-type validation; drill: malicious LLM output cannot reach DB.
Training-data poisoning	LLM03	Provenance log per dataset version; KAU + ABIDE-II hash-verified; drill: model trained on tampered data triggers eval-set failure.
Model denial-of-service	LLM04	Rate-limit + token-budget + max-context cap; per §41.4 sliding-window limits.
Supply-chain vulnerabilities	LLM05	Pinned model versions + SBOM + dependency scan (Trivy + pip-audit).
Sensitive info disclosure	LLM06	PII redaction at retrieval + audit log; drill: input PII never appears in output.
Insecure plugin / tool design	LLM07	Scope-grant policy engine (A3) + tool allowlist per agent; drill: scope-required tool denied without grant.
Excessive agency	LLM08	HITL gate (A10) on every non-routine action; drill: agent cannot bypass HITL.
Overreliance	LLM09	Plain-language uncertainty communication; caregiver-facing confidence band + “not a diagnosis” disclaimer.
Model theft	LLM10	API rate-limit + query-pattern anomaly detection + watermarked outputs.
Jailbreak / role-play bypass	beyond Top 10	Multi-layer guardrails (input + output + tool); council-vote dissent surfaces bypass attempts.
Hallucination	beyond Top 10	RAG citation grounding (A8) + Ragas faithfulness ≥ 0.85 gate; drill: uncited claim flagged.
Bias amplification	beyond Top 10	Disparate impact ≥ 0.80 + equal-opportunity gap $< 5\%$ pre-deploy gate; per §48 explainability.

Table 26: AI security threat catalogue (OWASP LLM Top 10 + 3 beyond) with RGAIG mitigation per row.

14.27 E27 — MLOps Lifecycle + 12-Tier Testing Matrix

RGAIG governs the full MLOps lifecycle (data → validate → feature → train → eval → registry → deploy → monitor → retrain) and a 12-tier testing matrix (unit through DDoS resilience) that gates every commit + every release. A model that does not pass every applicable tier is not deployable; the table below names each lifecycle step and each testing tier with its RGAIG gate.

#	MLOps step	Tool / framework	Gate + audit field
M1	Data validation	Great Expectations + Soda Core	Schema + quality + drift checks pre-train.
M2	Feature engineering	Feast (or in-house feature store)	Feature lineage + versioning.
M3	Model training	PyTorch + scikit-learn + Optuna	Reproducible random seed; HW + SW pinned.
M4	Evaluation (eval-set)	Held-out + Ragas (RAG) + DeepEval + Promptfoo (LLM)	Faithfulness ≥ 0.85 , accuracy gate per model card.
M5	Model registry	MLflow / DVC	Model + prompt + version + owner; rollback path.
M6	Deploy (canary / shadow)	Argo Rollouts + feature flags	Per §47.7 4-layer rollback.
M7	Monitoring	Prometheus + Evidently + Why-Labs OSS	Accuracy hourly + drift PSI daily + fairness weekly.
M8	Retrain triggers	Custom + scheduled cron	Accuracy drop $> 5\%$ OR drift PSI > 0.2 OR monthly.

12-tier testing matrix (per §64.30)

T1	Unit	pytest / vitest	$\geq 80\%$ line coverage; every commit (CI gate).
T2	Integration	pytest + real DB + real cache	Cross-service flows; every PR.
T3	API	httpx + schemathesis	Every endpoint $\times$ every status code.
T4	Boundary	hypothesis + faker	Min/max + off-by-one + empty + overflow.
T5	Process / drill	§43 drills with ≥ 3 negative assertions	E2E flow per process; commit-touched.
T6	Smoke	playwright / curl	Top 20 critical endpoints + UI routes; every deploy.
T7	Performance	k6 / locust	p95 latency per endpoint $< \text{SLA}$; nightly.
T8	Load (5-phase)	k6 / locust	Smoke + load + stress + soak + spike per §47.10.
T9	Security (SAST + DAST)	Semgrep + OWASP ZAP	Pre-release; auth env only.
T10	AI eval (Ragas + DeepEval + Promptfoo)	Per-prompt-version regression set	Faithfulness + answer-relevance + context-precision gates.
T11	Prompt-injection / red-team	Garak + Promptfoo + Rebuff	Pre-release; per LLM01–LLM10.
T12	DDoS resilience	wrk + chaos-mesh	Sustained 10 $\times$ peak; connection-flood survival.

Table 27: MLOps lifecycle (M1–M8) + 12-tier testing matrix (T1–T12) with RGAIG gates.

14.28 E28 — AI Discipline Map: 11 Cross-Cutting Concerns

RGAIG operationalises 11 cross-cutting AI disciplines that an examiner can audit independently: each row names the discipline, the RGAIG control + audit-row field, and the measurable threshold the dissertation defends. A model that lacks evidence for any row is non-deployable per the §47.11 pre-release gates.

#	AI discipline	RGAIG control + audit field	Measurable threshold + reporting
D1	PII protection	De-identification at source; AES-256 at rest; TLS 1.3 in transit; PII scanner (Presidio); audit-row pii_flag	0 PII in audit-row inputs; 0 leakage in outputs (per LLM06 drill).
D2	Consensus AI (Council voting)	3-stage author + reviewer + chair vote; audit-row author_proposal + reviewer_critique + chair_decision	Dissent surfaced as first-class signal; HITL escalation on disagreement (per A7).
D3	Responsible AI (RAI)	4-principle composite: fairness + accountability + transparency + privacy	Per §38 governance; pre-deploy + weekly drift checks.
D4	Explainable AI (XAI)	Two-layer: caregiver-readable plain language + clinician-grade SHAP / LIME / Integrated Gradients	SHAP top- <i>k</i> coverage $\geq 90\%$; caregiver-comprehension score Likert $\geq 5/7$.
D5	Interpretable AI (post-hoc surrogate)	Decision-tree surrogate depth ≤ 4 approximating black-box predictions; rule extraction	Surrogate fidelity ≥ 0.85 vs. black-box; rules.txt published per model card.
D6	Governance AI (RGAIG)	6-level maturity instrument + 24-control catalogue + multi-jurisdiction crosswalk	RGAIG maturity score ≥ 3 for deploy gate; crosswalk coverage $\geq 90\%$.
D7	Secure AI (OWASP LLM Top 10 + adversarial)	Per §E26 catalogue; Garak red-team + Promptfoo + Rebuff drills	Prompt injection rejection $\geq 95\%$; jailbreak survival $\geq 99\%$.
D8	Ethical AI (Mittelstadt + Belmont + EU AI Act Art. 86)	10-dimension safeguard table (privacy + pediatric + consent + beneficence + autonomy + fairness + hallucination + overtrust + AI disclosure + IP)	Per §3.9 ethics; pre-deploy gate + quarterly review.
D9	Fairness AI (statistical + counterfactual)	Disparate impact ≥ 0.80 + equal-opportunity gap $< 5\%$ + per-subgroup audit (IBM AIF360 + Fairlearn)	Pre-deploy gate + weekly drift; per-protected-attribute reporting.
D10	Output-First strategy	Ragas faithfulness + answer-relevance + context-precision; CI gate blocks merge on regression	Faithfulness ≥ 0.85 ; context-precision ≥ 0.75 ; answer-relevance ≥ 0.80 .
D11	Performance AI (latency + cost + throughput)	p95 latency + tokens-per-decision + cost-per-decision; OTel + Prometheus + Langfuse	p95 latency ≤ 2 s; cost-per-decision tracked per audit row; budget alerts at 80% / 100%.

Table 28: AI discipline map — 11 cross-cutting concerns with RGAIG control + measurable threshold.

14.29 E29 — NIST AI RMF and ISO/IEC 42001 Detailed Compliance Mapping

The two flagship AI-governance standards — NIST AI RMF 1.0 (2023) and ISO/IEC 42001:2023 — are mapped here at function / clause level to RGAIG controls. This is the procurement-grade evidence the hospital uses to

defend AI adoption to compliance + board + regulator.

E29a — NIST AI RMF 1.0 (2023) Mapping

NIST AI RMF organises trustworthy AI around four core functions (Govern, Map, Measure, Manage) with ~72 subcategories. The table below maps the four functions to RGAIG controls + audit-row evidence; a complete subcategory crosswalk lives in Appendix H.

Function	Purpose	RGAIG control + audit evidence
GOVERN	Risk management culture, policy, accountability, oversight	RGAIG 6-level maturity instrument; Doctoral Supervisor + Co-Supervisor + DBA Programme Director sign-off; AI Governance Board defined in Ch. 5 §5.6.
MAP	Context, intended use, risk classification, impact assessment	Ch. 1 §1.7 Context of Contributing Organisations; EU AI Act risk-tier tag per model; intended-use disclosure in Appendix I.
MEASURE	Quantitative + qualitative + technical evaluation metrics	Ch. 3 §3.10 Evaluation Metrics + Ch. 4 §4.7; per-model card with accuracy + fairness + explainability + drift scores.
MANAGE	Prioritise, respond, monitor, communicate AI risks	Decision-audit row schema (per §38.3); incident response runbook; HITL escalation ladder; weekly drift monitoring.
Cross-cutting	Trustworthy characteristics: valid + reliable + safe + secure + accountable + transparent + explainable + privacy-enhanced + fair	Per E28 D3–D11 discipline map; D1 PII control covers privacy-enhanced.

Table 29: NIST AI RMF 1.0 — Govern / Map / Measure / Manage + cross-cutting characteristics mapped to RGAIG.

E29b — ISO/IEC 42001:2023 (AI Management System) Mapping

ISO/IEC 42001:2023 is the first certifiable AI management system standard, structured like ISO 9001 / 27001 with Plan-Do-Check-Act discipline. The table below maps the 10 ISO clauses + Annex A controls to RGAIG controls + audit evidence; certifying body evidence is filed under Appendix D.

Clause	ISO/IEC 42001 area	RGAIG control + audit evidence
4	Context of the organisation	Ch. 1 §1.7 + Ch. 3 §3.4 (stakeholder + dataset context); ASD pediatric scope contract.
5	Leadership + AI policy	Sign-off chain: Supervisor + Co-Supervisor + DBA Programme Director; AI Governance Board (Ch. 5 §5.6).
6	Planning + AI objectives + risk treatment	RGAIG control catalogue (24-control); risk register (per §E13); AI objectives per O1–O6.
7	Support: resources + competence + awareness + communication + documented info	Training plan in Ch. 5 §5.4; documented info in Appendices A–J.
8	Operation: operational planning + AI system impact assessment + lifecycle	MLOps lifecycle per E27 (M1–M8); impact assessment per model card; operational HITL gate.
9	Performance evaluation: monitoring + measurement + analysis + management review	Per E28 D6 Governance AI thresholds; weekly drift + monthly governance review.
10	Improvement: nonconformity + corrective action + continual improvement	Incident postmortem → control update → re-drill; PDCA cycle.
Annex A	Controls (A.1–A.10): policies + objectives + data + AI system + lifecycle + third-party + impact + use + reporting	Mapped row-by-row in Appendix H; covers all RGAIG layers (data + AI + governance + HITL + value + audit).
Annex B	Implementation guidance	Per E16 RGAIG Apply 7-step playbook.
Annex C	Risk sources + categories	Per §E13 Risk Register + per E26 OWASP LLM Top 10.
Annex D	Use cases across sectors	ASD pediatric use-case fully documented; cross-sector extension is Phase 2 future work.

Table 30: ISO/IEC 42001:2023 clauses + annexes mapped to RGAIG controls + audit evidence.

14.30 E30 — SMART Test of Research Aim and Objectives

The research aim and the six objectives (O1–O6) are evaluated against the SMART criteria (Specific, Measurable, Achievable, Relevant, Time-bound); each row certifies the SMART qualification per objective. An objective failing any SMART dimension is reformulated until it passes; the table below is the certified version submitted with the dissertation.

E30a — SMART Aim Statement (Reformulated for SMART Compliance)

The research aim is restated below in SMART form so that each criterion is explicitly satisfied; the table immediately after maps each SMART dimension to its evidence.

SMART dimension	Aim-statement evidence
Aim (SMART form)	<i>“By June 2026, develop and empirically validate the Responsible Generative AI Governance (RGAIG) framework — a 5-layer + transverse-audit governance fabric for B2C-primary, clinician-supported AI-assisted retrospective EEG-based pediatric ASD screening support — evaluated via PLS-SEM (n=131 clinicians, $p < 0.05$ on H1–H5) + qualitative expert panel (n=12) + retrospective benchmark on KAU + ABIDE-II datasets, achieving RGAIG maturity score ≥ 3 of 6 and multi-jurisdiction crosswalk coverage $\geq 90\%$.”</i>
S — Specific	Names the artefact (RGAIG framework), the disorder (ASD), the modality (EEG), the stakeholder positioning (B2C-primary), the safety layer (clinician-supported), and the scope boundary (retrospective).
M — Measurable	PLS-SEM p-value < 0.05 ; RGAIG maturity score 0–6 with deploy gate ≥ 3 ; multi-jurisdiction crosswalk coverage %; faithfulness ≥ 0.85 ; HITL overhead $< 10\%$ clinician time.
A — Achievable	131-clinician PLS-SEM + 12-expert panel already conducted; KAU + ABIDE-II datasets accessible; candidate’s agenticfinder reference codebase (305 GB; 198 trained models) supplies the technical anchor.
R — Relevant	ASD prevalence rising globally; pediatric screening burden documented; EU AI Act + ISO/IEC 42001 + NIST AI RMF live in 2024–2026 creating the governance window.
T — Time-bound	June 2026 submission target with quarterly milestones (per §E15 timeline); Phase 2 Indian cohort explicitly bounded as post-submission.

Table 31: SMART aim — reformulated statement + per-criterion evidence.

E30b — SMART Test of Objectives O1–O6

Each of the six objectives is scored against the five SMART criteria; a green tick certifies the criterion is satisfied with named evidence, a yellow mark flags improvement, a red cross blocks submission until corrected. All six objectives pass all five SMART criteria in the certified version below.

#	Objective	S	M	A	R	T	Evidence anchor
O1	Investigate governance gaps in pediatric AI screening support	✓	✓	✓	✓	✓	Ch. 1 + Ch. 2 + Q1/Q4 2025 milestone (E15 timeline).
O2	Evaluate existing frameworks (NIST AI RMF, ISO/IEC 42001, EU AI Act, DPDP)	✓	✓	✓	✓	✓	E29 NIST + ISO mapping; crosswalk coverage $\geq 90\%$ gate.
O3	Design the RGAIG framework artefact (5 layers + audit fabric)	✓	✓	✓	✓	✓	Figure E3 + Appendix G; Q4 2025 design milestone.
O4	Validate framework effectiveness empirically (PLS-SEM + panel + benchmark)	✓	✓	✓	✓	✓	H1–H5 supported at $p < 0.05$; n=131 + n=12 + KAU + ABIDE-II.
O5	Assess business + operational + ROI implications	✓	✓	✓	✓	✓	Ch. 5 §5.2 5-dimension implications; 70% time-to-support reduction (illustrative).
O6	Articulate adoption pathway + maturity instrument	✓	✓	✓	✓	✓	6-level RGAIG maturity (E16 RGAIG Apply); 24-control catalogue.

Table 32: SMART test of objectives O1–O6 — all six pass all five criteria with named evidence.

14.31 E31 — C4 Architecture Model (L1 System Context + L2 Containers + L3 Components + L4 Code)

The Simon Brown C4 model provides four hierarchical architecture views; each level zooms in one step and is the canonical examiner-facing artefact per §47 Architecture standards. The table + figure below document L1 + L2; L3 + L4 are referenced (full diagrams in Appendix I).

E31a — L1 System Context (Actors + External Systems)

L1 names every actor and external system that surrounds RGAIG; the table below catalogues 4 actor types + 3 external system types + the RGAIG boundary.

Element type	Element	Interaction with RGAIG
Actor (primary)	Caregiver / Family	Submits query; receives plain-language explanation + audit row on demand.
Actor (beneficiary)	Child (pediatric patient)	Indirect beneficiary; protected by pediatric safeguards (ICMR + DPDP).
Actor (safety)	Clinician / Neuropsychiatrist	HITL gate; reviews + approves / escalates / overrides; retains diagnostic authority.
Actor (governance)	Hospital Admin / Compliance Officer	Procurement decision via RGAIG maturity instrument; audit-row evidence to board + regulator.
External system	Reference datasets (KAU + ABIDE-II)	Read-only data ingress; provenance-logged.
External system	IRB / Regulator (EU AI Act, FDA SaMD, DPDP)	Compliance evidence egress; signed batch audit-row export.
External system	Audit Registry (7-yr retention)	Long-term audit-row archive; tamper-evident hash chain.
RGAIG boundary	RGAIG Framework	5 internal layers + transverse audit fabric (per Figure E3).

Table 33: C4 L1 System Context — 4 actor types + 3 external system types + RGAIG boundary.

E31b — L2 Containers (Deployable Services Inside RGAIG)

L2 decomposes the RGAIG boundary into 8 containers (deployable services) + 4 stores; each container has a tech stack + a single responsibility.

#	Container	Tech stack	Responsibility
C1	Caregiver UI (B2C)	Next.js + GraphQL	B2C-primary entry point; plain-language XAI display.
C2	Clinician Portal (B2B)	Next.js + RBAC	HITL review console + scope-grant log.
C3	API Gateway	FastAPI + Kong	Auth + rate-limit + telemetry + tenant routing.
C4	Agentic Orchestrator	Python + Temporal	Planner + decomposer + policy engine; per E23 A1–A3.
C5	RAG Engine	LangChain + Perplexity-style citation	A8 + A9 hybrid retrieval; ChromaDB vector store.
C6	ML Core (model registry + inference)	PyTorch + MLflow + SHAP	198-model registry; SHAP / LIME explanation.
C7	Council of Agents	3 Ollama instances (author + reviewer + chair)	A7 voting pattern; dissent surfaces uncertainty.
C8	HITL Service	FastAPI + Web-Socket	Clinician review + scope-grant + escalation router.
S1	Audit Store	PostgreSQL + Kafka events	Decision-audit rows; 7-yr retention; tamper-evident.
S2	Vector Store	ChromaDB (1.4 GB embeddings)	Per A9 hybrid retrieval.
S3	Graph + RDF Store	Neo4j + Apache Jena	RGAIG ontology + multi-hop reasoning (per E25).
S4	Time-series Store	TimescaleDB	EEG signal archive + SLI metrics for drift detection.

Table 34: C4 L2 Containers — 8 deployable services + 4 persistent stores inside RGAIG boundary.

E31c — L3 Components + L4 Code (Reference Pointers)

L3 zooms into one container (e.g., Agentic Orchestrator → Planner + Decomposer + Policy + Reflection components); L4 shows class / module structure of a single component. Both levels are detailed in Appendix I (architecture diagrams) of the main thesis; this v3 doc references them.




Figure 9: Agentic orchestration pattern — user → planner → decomposer → policy → MCP → Perplexity-RAG → council → HITL.

Level	Zoom target	Where documented
L3	Components inside Agentic Orchestrator (C4): Planner + Decomposer + Policy + Reflection + Memory	Main thesis Appendix I (architecture diagrams).
L3	Components inside RAG Engine (C5): Indexer + Retriever + Reranker + Generator + Citation-tracker	Main thesis Appendix I.
L3	Components inside Council (C7): Author-agent + Reviewer-agent + Chair-agent + Audit-row writer	Main thesis Appendix I.
L4	Class diagrams for: Planner / RetrievalAdapter / VotingChair / HITL-Router / Decision-AuditRow	Reference repo <code>agenticfinder</code> ; UML diagrams in main thesis Appendix I.

Table 35: C4 L3 Components + L4 Code — reference pointers to main thesis Appendix I.

E23b — Agentic Orchestration Pattern (Flowchart)

The orchestration flowchart shows how a caregiver query traverses the agentic stack: user → planner → decomposer → policy → MCP-mediated tool calls → council vote → HITL → delivery. Every step writes a decision-audit row; the council vote is the deterministic-disagreement gate that surfaces uncertainty as a first-class signal.

End of DBA Dissertation v3 (GGU-Strict Structure with Enhanced Additions). Coverage matrix verified at top: all 47 GGU-required topics + 6 enhanced front-matter / appendix items + 22 enhanced additions (6 figures + 17 tables) all addressed. Pipeline documented end-to-end (E21 23-step flow). Stakeholder decomposition B2B + B2C explicit (E22). Bibliography aligned with main thesis. Title aligned with main thesis. Supervisor (C. Ranjeeth Kumar) + Co-Supervisor (Sumitra Padmanabhan) named.